

# EECS E6690: SL for Bio & Info

## Lecture 8: Support Vector Machines, Optimization, and Gene Expression Classification

Prof. Predrag R. Jelenković

Time: Tuesday 4:10-6:40pm

Dept. of Electrical Engineering  
Columbia University , NY 10027, USA

Office: 812 Schapiro Research Bldg.

Phone: (212) 854-8174

Email: [predrag@ee.columbia.edu](mailto:predrag@ee.columbia.edu)

URL: <http://www.ee.columbia.edu/~predrag>

# Final Project Outline

- ▶ Done in groups of 4 students - assemble the groups
- ▶ Deliverables: 15+ page **paper** & **presentation** with slides
- ▶ **Due:** during the finals week: Dec 16 - 23, very likely **Dec 17**.  
One slot for presentations on **Tue, Dec 14, 4:10-6:40pm**.
- ▶ **Data Repositories:** First, select a paper(s) from either:
  - ▶ UC Irvine Machine Learning Repository  
<https://archive.ics.uci.edu/ml/datasets.php>
  - ▶ GEO Data Repository <https://www.ncbi.nlm.nih.gov/geo/>,  
or Bioconductor Datasets  
<http://www.bioconductor.org/packages/release/data/experiment/>
- ▶ **Final Paper Outline:** 5 sections
  1. **Introduction:** e.g., describe the application area, problems considered, etc
  2. **Data set(s) and paper(s):** e.g., describe data in detail, what was done in the paper(s), common stat/machine learning tools, etc
  3. **Reproduce the results from the paper(s)**
  4. **Try different techniques learned in class, or propose new ones**
  5. **Discussion and conclusion:** e.g., compare different techniques, pros and cons, future work, etc

# Bioconductor and Additional Datasets

- ▶ Bioconductor provides tools in R for the analysis genomic data:  
[https://https://www.bioconductor.org/](https://www.bioconductor.org/)

- ▶ Installing Bioconductor:

[https://https://www.bioconductor.org/install/](https://www.bioconductor.org/install/)

Run the following code:

```
if (!requireNamespace("BiocManager", quietly = TRUE))  
  install.packages("BiocManager")  
BiocManager::install(version = "3.12")
```

- ▶ Then, install Bioconductor packages:

[https://www.bioconductor.org/install/](https://www.bioconductor.org/install/#install-bioconductor-packages)  
#install-bioconductor-packages

- ▶ Datasets supported by Bioconductor: <http://www.bioconductor.org/packages/release/data/experiment/>

# Last lecture: Bootstrap Methods

- ▶ Stochastic search
- ▶ Works for both: classifiers or regressions
- ▶ Bumping, Bagging, Random forests, Boosting
- ▶ Train a classifier or regression model  $\hat{f}_0$  on  $\mathcal{Z}$
- ▶ For  $b = 1, \dots, B$ :
  1. Draw a bootstrap sample  $\mathcal{Z}^{*b}$  of size  $n$  from training data
  2. Train a classifier or regression model  $\hat{f}_b$  on  $\mathcal{Z}^{*b}$
- ▶ **Bumping**: Select the best model, e.g.,

$$\hat{b} = \arg \min_{0 \leq b \leq B} \sum_{i=1}^n \left( y_i - \hat{f}_b(z_i) \right)^2$$

- ▶ **Bagging**: average out - reduces variance
- For a “new” point  $\mathbf{x}_0$ , compute:

$$\hat{f}_{\text{avg}}(\mathbf{x}_0) = \frac{1}{B} \sum_{b=1}^B \hat{f}_b(\mathbf{x}_0)$$

# Last lecture: Random forests

Works for both: classifiers or regressions

- ▶ Improvement over bagged trees
- ▶ Idea: Decorrelate trees
  - ▶ Still learn a tree on each bootstrap set
  - ▶ To split a region, consider only a subset of predictors
- ▶ Input parameter:  $m \leq p$ , often  $m \approx \sqrt{p}$
- ▶ For  $b = 1, \dots, B$ 
  - ▶ Draw a bootstrap sample  $\mathbf{Z}^{*b}$  of size  $n$  from the training data
  - ▶ Train a tree classifier on  $\mathbf{Z}^{*b}$ , each split is computed as:
    - ▶ Randomly select  $m$  predictors, newly chosen for each  $b$
    - ▶ Make the best split restricted to that subsets of predictors

# Last lecture: Boosting for regression

Slow learning

1. Set  $\hat{f}(\mathbf{x}_i) = 0$  and  $r_i = y_i$  for all  $i$  in the training set
2. For  $b = 1, \dots, B$ , repeat:
  - 2.1 Fit a tree  $\hat{f}_b$  with  $d$  splits ( $d + 1$  terminal nodes) to training data  $(\mathbf{X}, \mathbf{r})$
  - 2.2 Update  $\hat{f}$  by adding in a shrunk version of the new tree:

$$\hat{f}(\mathbf{x}) \leftarrow \hat{f}(\mathbf{x}) + \lambda \hat{f}_b(\mathbf{x})$$

- 2.3 Update residuals:

$$r_i \leftarrow r_i - \lambda \hat{f}_b(\mathbf{x}_i)$$

3. Output the boosted model:

$$\hat{f}(\mathbf{x}) = \sum_{b=1}^B \lambda \hat{f}_b(\mathbf{x})$$

► Notes:

- $\lambda$  is a small positive number (e.g., 0.01 or 0.001)
- Often  $d = 1$  works

# Boosting for Classification: AdaBoost

- ▶ Due to Freund and Schapire (1997)
- ▶ Consider two classes:  $Y \in \{-1, 1\}$
- ▶ For a classifier  $G(x) \in \{-1, 1\}$ , the training error is

$$e_m = \frac{1}{n} \sum_{i=1}^n 1_{\{y_i \neq G(x_i)\}}$$

- ▶ Main idea: construct **weak classifiers**
  - ▶ Weak classifier: slightly better than random guessing (50% error)
  - ▶ Sequentially, construct weak classifiers,  $G_m(x)$ , on modified training data.
- ▶ Final classifier, combination of weak classifiers through a weighted majority vote

$$G(x) = \text{sign} \left[ \sum_{m=1}^M \alpha_m G_m(x) \right]$$

# Boosting for Classification: AdaBoost

1. Set  $w_i = 1/n, i = 1, 2, \dots, n$ , where  $n$  is the number of training points.
2. For  $m = 1, \dots, M$ , repeat:
  - (a) Fit a (weak) classifier  $G_m(x)$  to the training data using weights  $w_i$ .
  - (b) Compute the weighted error

$$e_m = \frac{\sum_{i=1}^n w_i 1_{\{y_i \neq G_m(x_i)\}}}{\sum_{i=1}^n w_i}$$

- (c) Compute

$$\alpha_m = \log((1 - e_m)/e_m).$$

- (d) Update

$$w_i \leftarrow w_i \exp(\alpha_m 1_{\{y_i \neq G_m(x_i)\}}), \quad i = 1, 2, \dots, n.$$

3. Final classifier


$$G(x) = \text{sign} \left[ \sum_{m=1}^M \alpha_m G_m(x) \right]$$



# Last lecture: The Maximal Margin Classifier

## Optimal hyperplane

- ▶ *Margin*: Distance from an observation to the hyperplane
- ▶ *Maximal margin hyperplane*: One whose smallest margin is maximal
- ▶ *Support vectors*: points that support the maximal margin hyperplane




# Problems

## Non-separable case

- ▶ There is no hyperplane that separates the two classes

## Highly sensitive to support vectors

- ▶ The hyperplane moves if one moves or introduces new support vector points



Need a "softer" separator

# Last lecture: Support Vector Classifier

- ▶ Greater robustness to individual observations
- ▶ More general: Works for most points

The hyperplane is the solution to

$$\max_{\beta_j, \epsilon_j} M \quad (1)$$

$$\text{subject to } \sum_1^p \beta_j^2 = 1 \quad (2)$$

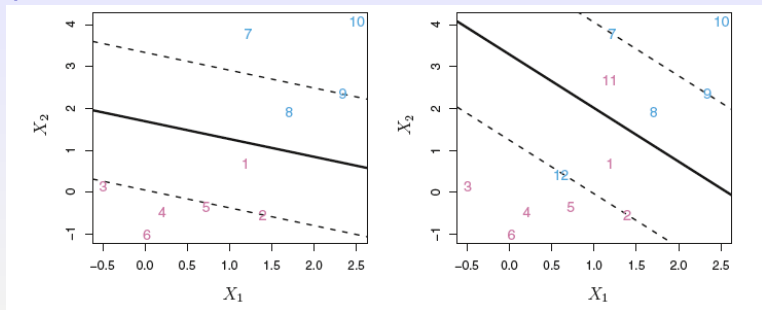
$$y_i(\beta_0 + \beta_1 x_{i1} + \beta_2 x_{i2} + \cdots + \beta_p x_{ip}) \geq M(1 - \epsilon_i), \quad \forall i \quad (3)$$

$$\epsilon_i \geq 0, \sum_{i=1}^n \epsilon_i \leq C \quad (4)$$

Notes:

- ▶  $\epsilon_i$  - *slack variables*
  - ▶  $\epsilon_i = 0$  -  $i$ -th observation on the correct side of the margin
  - ▶  $0 < \epsilon_i < 1$  -  $i$ -th observation on the wrong side of the margin
  - ▶  $\epsilon_i > 1$  -  $i$ -th observation on the wrong side of the hyperplane
- ▶  $C$  - budget for slackness

# Example



- ▶ Left:
  - ▶ Purple observations: 3,4,5, and 6 = correct side of the margin; 2 is on the margin, and 1 is on the wrong side of the margin.
  - ▶ Blue observations: 7 and 10 = correct side of the margin; 9 is on the margin, and 8 is on the wrong side of the margin.
- ▶ Right: Same as left panel with two additional points, 11 and 12. 11 and 12 = wrong side of both the hyperplane and the margin.
- ▶ **Robustness**: the hyperplane on the right did not move much (!)

# General Support Vector Machines

Nonlinear decision boundary - example

- ▶  $p$  features:  $X_1, X_2, \dots, X_p$
- ▶ expand bases -  $2p$  features:  $X_1, X_1^2, X_2, X_2^2, \dots, X_p, X_p^2$

Fit the hyperplane to the expanded bases

$$\max_{\beta_j, \epsilon_j} M \quad (5)$$

$$\text{subject to } \sum_1^p \sum_{k=1}^2 \beta_{jk}^2 = 1 \quad (6)$$

$$y_i(\beta_0 + \sum_{j=1}^p \beta_{j1}x_{ij} + \sum_{j=1}^p \beta_{j2}x_{ij}^2) \geq M(1 - \epsilon_i), \quad \forall i \quad (7)$$

$$\epsilon_i \geq 0, \sum_{i=1}^n \epsilon_i \leq C \quad (8)$$

- ▶ Nonlinear surface since these are quadratic expressions  
Quadratic can be replaced by polynomial of any degree
- ▶ The SVM explores further this idea using kernels

# Understanding geometry: Distance from a hyperplane

Consider a hyperplane in  $p$ -dimensions

$$\beta_0 + \langle \boldsymbol{\beta}, \mathbf{x} \rangle = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \cdots + \beta_p x_p = 0$$

where  $\boldsymbol{\beta} = (\beta_1, \dots, \beta_p)$  and  $\mathbf{x} = (x_1, \dots, x_p)$

- **Claim:**  $\boldsymbol{\beta}$  is a perpendicular vector to this hyperplane

**Proof:** Let  $\mathbf{x}$  and  $\mathbf{x}'$  be two points on this hyperplane. Then

$$\beta_0 + \langle \boldsymbol{\beta}, \mathbf{x} \rangle - (\beta_0 + \langle \boldsymbol{\beta}, \mathbf{x}' \rangle) = \langle \boldsymbol{\beta}, (\mathbf{x}' - \mathbf{x}) \rangle = 0$$

- Perpendicular unit vector

$$\frac{\boldsymbol{\beta}}{\|\boldsymbol{\beta}\|}$$

where  $\|\boldsymbol{\beta}\| \equiv \|\boldsymbol{\beta}\|_2$  is the usual euclidian norm.

- **Signed distance** to the hyperplane: Let  $\mathbf{x}_0$  be a point outside the hyperplane and  $\mathbf{x}$  inside (note  $\langle \boldsymbol{\beta}, \mathbf{x} \rangle = -\beta_0$ )

$$\frac{\langle \boldsymbol{\beta}, (\mathbf{x}_0 - \mathbf{x}) \rangle}{\|\boldsymbol{\beta}\|} = \frac{\langle \boldsymbol{\beta}, \mathbf{x}_0 \rangle + \beta_0}{\|\boldsymbol{\beta}\|}$$

# SVC Geometry

Using the preceding distance formula, we get

$$\begin{aligned} & \max_{\beta_j, \epsilon_j} M \\ \text{subject to } & y_i \frac{\langle \beta, \mathbf{x} \rangle + \beta_0}{\|\beta\|} \geq M(1 - \epsilon_i), \quad \forall i \\ & \epsilon_i \geq 0, \sum_{i=1}^n \epsilon_i \leq C \end{aligned}$$

Recall that in the last lecture we set  $\|\beta\|^2 = \sum_1^p \beta_j^2 = 1$ , in which case  $\beta$  is the unit normal vector.

- We can set  $M = 1/\|\beta\|$  in the preceding optimization

**Reason:** If  $(\beta_0, \beta)$  satisfies the preceding equations, then any scaled version of it satisfies, and in particular, we can set  $M = 1/\|\beta\|$ , instead of  $\|\beta\| = 1$ .

# SVC: Convex Optimization

Next, optimizing  $\max(1/\|\beta\|)$  is equivalent to  $\min(\|\beta\|^2/2)$

$$\begin{aligned} & \min_{\beta_j, \epsilon_j} \frac{\|\beta\|^2}{2} \\ \text{subject to } & y_i(\langle \beta, \mathbf{x} \rangle + \beta_0) \geq (1 - \epsilon_i), \quad \forall i \\ & \epsilon_i \geq 0, \sum_{i=1}^n \epsilon_i \leq C \end{aligned}$$



# Constrained Optimization: Crash Course

## Primal problem

$$\begin{aligned} \min_{\mathbf{x} \in R^n} f(\mathbf{x}) \\ \text{subject to } f_i(\mathbf{x}) \leq 0, \quad i = 1, \dots, m \\ h_i(\mathbf{x}) = 0, \quad i = 1, \dots, p \end{aligned} \tag{9}$$

## Lagrangian

- Incorporate the constraints into one equation

$$L(\mathbf{x}, \boldsymbol{\lambda}, \boldsymbol{\mu}) := f(\mathbf{x}) + \sum_{i=1}^m \lambda_i f_i(\mathbf{x}) + \sum_{i=1}^p \mu_i h_i(\mathbf{x}) \tag{10}$$

*Lagrange multiplier vectors or dual vectors*

$$\boldsymbol{\lambda} = (\lambda_1, \dots, \lambda_m), \quad \boldsymbol{\mu} = (\mu_1, \dots, \mu_p)$$

# Constrained Optimization

## Lagrange dual function

$$\begin{aligned} g(\boldsymbol{\lambda}, \boldsymbol{\mu}) &:= \inf_{\boldsymbol{x} \in R^n} L(\boldsymbol{x}, \boldsymbol{\lambda}, \boldsymbol{\mu}) \\ &= \inf_{\boldsymbol{x} \in R^n} \left( f(\boldsymbol{x}) + \sum_{i=1}^m \lambda_i f_i(\boldsymbol{x}) + \sum_{i=1}^p \mu_i h_i(\boldsymbol{x}) \right) \end{aligned} \quad (11)$$

- ▶ Let  $f^*$  = optimal value of the primal problem
- ▶ **Lower bound:** for any  $\boldsymbol{\lambda} \geq 0$

$$f^* \geq g(\boldsymbol{\lambda}, \boldsymbol{\mu})$$

since, for any feasible point  $\tilde{\boldsymbol{x}}$  (  $\tilde{\boldsymbol{x}}$  satisfies the primal problem)

$$\sum_{i=1}^m \lambda_i f_i(\tilde{\boldsymbol{x}}) + \sum_{i=1}^p \mu_i h_i(\tilde{\boldsymbol{x}}) = \sum_{i=1}^m \lambda_i f_i(\tilde{\boldsymbol{x}}) \leq 0$$

since all  $\lambda_i \geq 0$ .

# Constrained Optimization

## Lagrange dual problem

$$\begin{aligned} & \max g(\boldsymbol{\lambda}, \boldsymbol{\mu}) \\ \text{subject to } & \boldsymbol{\lambda} \geq 0 \end{aligned} \tag{12}$$

- ▶ Let  $g^* = g(\boldsymbol{\lambda}^*, \boldsymbol{\mu}^*)$  be the optimal value of the dual problem  
 $\boldsymbol{\lambda}^*, \boldsymbol{\mu}^*$  are the optimal Lagrange multipliers

- ▶ Clearly

$$f^* \geq g^*$$

- ▶ When are primal and dual problems equal, i.e.,  $f^* = g^*$ ?
- ▶ **Duality gap**

$$f(\boldsymbol{x}) - g(\boldsymbol{\lambda}, \boldsymbol{\mu})$$

- ▶ Can be used for computation as a **stopping criterion**.

# Constrained Optimization

## Complementary slackness

- If primal and dual problem values are equal,  $f^* = g^*$ , and attained at values  $f^* = f(\mathbf{x}^*)$  and  $g^* = g(\boldsymbol{\lambda}^*, \boldsymbol{\mu}^*)$ , then

$$\lambda_i^* f_i(\mathbf{x}^*) = 0, \quad i = 1, \dots, m$$

- **Proof**

$$\begin{aligned} f^* &= f(\mathbf{x}^*) = g^* = g(\boldsymbol{\lambda}^*, \boldsymbol{\mu}^*) \\ &= \inf_{\mathbf{x} \in \mathbb{R}^n} \left( f(\mathbf{x}) + \sum_{i=1}^m \lambda_i^* f_i(\mathbf{x}) + \sum_{i=1}^p \mu_i h_i^*(\mathbf{x}) \right) \\ &\leq f(\mathbf{x}^*) + \sum_{i=1}^m \lambda_i^* f_i(\mathbf{x}^*) + \sum_{i=1}^p \mu_i h_i^*(\mathbf{x}^*) \\ &\leq f(\mathbf{x}^*) \end{aligned}$$

since  $\lambda_i^* \geq 0$ ,  $f_i(\mathbf{x}^*) \leq 0$  and  $h_i^*(\mathbf{x}^*) = 0$ .

Hence,  $\sum_{i=1}^m \lambda_i^* f_i(\mathbf{x}^*) = 0$ , and therefore  $\lambda_i^* f_i(\mathbf{x}^*) = 0$ .

# Karush-Kuhn-Tucker (KKT) Conditions

## Necessary conditions for nonconvex problems

- ▶ Let  $\mathbf{x}^*$  and  $\boldsymbol{\lambda}^*, \boldsymbol{\mu}^*$  be any primal and dual points with zero duality gap.
- ▶ Then, the following KKT conditions hold

$$f_i(\mathbf{x}^*) \leq 0, \quad (\text{primal feasibility})$$

$$h_i(\mathbf{x}^*) = 0, \quad (\text{primal feasibility})$$

$$\lambda_i^* \geq 0, \quad (\text{dual feasibility})$$

$$\lambda_i^* f_i(\mathbf{x}^*) = 0, \quad (\text{complem. slackness})$$

$$\nabla f(\mathbf{x}^*) + \sum_{i=1}^m \lambda_i^* \nabla f_i(\mathbf{x}^*) + \sum_{i=1}^p \mu_i^* \nabla h_i(\mathbf{x}^*) = 0, \quad (\text{stationarity})$$

Recall  $\nabla f(\mathbf{x}) = \left( \frac{\partial f(\mathbf{x})}{\partial x_1}, \dots, \frac{\partial f(\mathbf{x})}{\partial x_n} \right)$  is the gradient

- ▶ 4th condition = complementary slackness  $\Rightarrow$   
 $\lambda_i = 0, f_i(\mathbf{x}^*) < 0$ :  $\mathbf{x}^*$  is **inside**, or  
 $\lambda_i > 0, f_i(\mathbf{x}^*) = 0$ :  $\mathbf{x}^*$  on the **boundary**

# Karush-Kuhn-Tucker (KKT) Conditions

## Sufficient conditions for convex problems

- ▶ Assume  $f, f_i$  are convex, and  $h_i$  are affine (linear)
- ▶ Let  $\tilde{x}$  and  $(\tilde{\lambda}, \tilde{\mu})$  be any points that satisfy KKT conditions
- ▶ Then,  $\tilde{x}$  and  $(\tilde{\lambda}, \tilde{\mu})$  are primal and dual optimal with zero duality gap

▶ **Proof:**

$$\begin{aligned} g(\tilde{\lambda}, \tilde{\mu}) &= L(\tilde{x}, \tilde{\lambda}, \tilde{\mu}) \\ &= f(\tilde{x}) + \sum_{i=1}^m \tilde{\lambda}_i f_i(\tilde{x}) + \sum_{i=1}^p \tilde{\mu}_i h_i(\tilde{x}) \\ &= f(\tilde{x}) \end{aligned}$$

We used in the second equality that  $L(x, \tilde{\lambda}, \tilde{\mu})$  is convex in  $x$  since  $\tilde{\lambda}_i \geq 0$  and  $h_i$  are affine, implying, by the last KKT condition, that  $\tilde{x}$  minimizes  $L(x, \tilde{\lambda}, \tilde{\mu})$ .

In the last line, we used  $h_i(x^*) = 0$  and  $\lambda_i^* f_i(x^*) = 0$  (complementary slackness)

# SVC Continued: Dual Problem

Incorporate the SVC constraints into the Lagrange function

$$L(\boldsymbol{\beta}, \beta_0, \boldsymbol{\alpha}, \boldsymbol{\mu}) = \frac{\|\boldsymbol{\beta}\|^2}{2} + c \sum_{i=1}^n \epsilon_i - \sum_{i=1}^n \alpha_i [y_i (\langle \boldsymbol{\beta}, \mathbf{x}_i \rangle + \beta_0) - (1 - \epsilon_i)] - \sum_{i=1}^n \mu_i \epsilon_i$$

Next, we define the **dual function**

$$g(\boldsymbol{\alpha}, \boldsymbol{\mu}) = \inf_{\boldsymbol{\beta}, \beta_0, \epsilon_i} L(\boldsymbol{\beta}, \beta_0, c, \boldsymbol{\alpha}, \boldsymbol{\mu})$$

Then, since  $L(\boldsymbol{\beta}, \beta_0, \boldsymbol{\alpha}, \boldsymbol{\mu})$  is quadratic in  $\beta_i$ , we can explicitly compute its derivatives with respect to  $\beta_i$  and  $\epsilon_i$

$$\boldsymbol{\beta} = \sum_{i=1}^n \alpha_i y_i \mathbf{x}_i \quad (\partial / \partial \boldsymbol{\beta})$$

$$0 = \sum_{i=1}^n \alpha_i y_i \quad (\partial / \partial \beta_0)$$

$$\alpha_i = c - \mu_i \quad (\partial / \partial \epsilon_i)$$

Plugging the preceding derivatives into  $L$ , yields an **explicit formula for  $g$** , after some algebra.

# SVC: Dual Problem

Now, dual problem is to maximize  $g(\boldsymbol{\alpha}, \boldsymbol{\mu})$ , i.e.,

$$\max_{\alpha_i} \left( \sum_{i=1}^n \alpha_i - \frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n \alpha_i \alpha_j y_i y_j \langle \mathbf{x}_i, \mathbf{x}_j \rangle \right)$$

$$\text{subject to } 0 \leq \alpha_i \leq c, \sum \alpha_i y_i = 0$$

KKT conditions to the rescue

$$\alpha_i [y_i (\langle \boldsymbol{\beta}, \mathbf{x}_i \rangle + \beta_0) - (1 - \epsilon_i)] = 0 \quad \text{slackness}$$

$$\mu_i \epsilon_i = 0$$

$$y_i (\langle \boldsymbol{\beta}, \mathbf{x}_i \rangle + \beta_0) - (1 - \epsilon_i) \geq 0$$

plus the preceding derivative equations.

From **slackness**,  $\alpha_i > 0$  only for support vectors, when

$$y_i (\langle \boldsymbol{\beta}, \mathbf{x}_i \rangle + \beta_0) = (1 - \epsilon_i)$$



# SVC Simplification

From the derivative condition (2 slides ago)

$$\begin{aligned}\hat{\beta} &= \sum_{i=1}^n \alpha_i y_i \mathbf{x}_i \\ &= \sum_{i \in \mathcal{S}} \alpha_i y_i \mathbf{x}_i\end{aligned}$$

And, thus, the classification hyperplane is easy to compute

$$f(x) = \beta_0 + \sum_{i \in \mathcal{S}} \alpha_i \langle \mathbf{x}_i, \mathbf{x} \rangle$$

# Support Vector Machines: Hilbert spaces

**Hilbert spaces:** Generalized linear spaces

- ▶ Let  $\phi(\mathbf{x}_i)$  be a transformation of feature variables such that
- ▶ **Kernel** -  $K(\mathbf{x}_i, \mathbf{x}_j)$  is positive definite
- ▶ **Generalized inner product:**

$$\langle \phi(\mathbf{x}_i), \phi(\mathbf{x}_j) \rangle = K(\mathbf{x}_i, \mathbf{x}_j)$$

- ▶ Hence, with arbitrary nonlinear transformation, Kernel, we can repeat the preceding optimization, and derive an SVC.

The resulting classifier is known as **SVM**, with the decision function taking the following nonlinear form in general

$$f(x) = \beta_0 + \sum_{i \in \mathcal{S}} \alpha_i K(x, x_i)$$

# Support Vector Machines: Kernels

## Kernels:

- ▶ *Linear*

$$K(x_i, x_j) = \langle x_j, x_i \rangle$$

- ▶ *Polynomial* - for positive integer  $d$

$$K(x_i, x_j) = (1 + \langle x_j, x_i \rangle)^d$$

- ▶ *Radial* - for  $\gamma > 0$  - (or Gaussian  $\gamma = 1/(2\sigma^2)$ )

$$K(x_i, x_j) = \exp \left( -\gamma \sum_{k=1}^p (x_{ik} - x_{jk})^2 \right)$$

# Deriving The Quadratic Kernel

## Kernels:

- ▶ Consider data with two predictors:  $x_i = (x_{i1}, x_{i2})$
- ▶ *Quadratic Kernel* ( $d = 2$ ):  $K(x_i, x_j) = (1 + \langle x_j, x_i \rangle)^2$
- ▶ We can obtain the preceding kernel by considering feature map

$$\phi(x_i) = (1, \sqrt{2}x_{i1}, \sqrt{2}x_{i2}, x_{i1}^2, \sqrt{2}x_{i1}x_{i2}, x_{i2}^2)$$

- ▶ Then, the inner product

$$\begin{aligned}\langle \phi(x_i), \phi(x_j) \rangle &= 1 + 2x_{i1}x_{j1} + 2x_{i1}x_{j1} + x_{i1}^2x_{j1}^2 \\ &\quad + 2x_{i1}x_{i2}x_{j1}x_{j2} + x_{i2}^2x_{j2}^2 \\ &= 1 + 2\langle x_i, x_j \rangle + \langle x_i, x_j \rangle^2 \\ &= (1 + \langle x_j, x_i \rangle)^2 = K(x_i, x_j)\end{aligned}$$

# Finding Feature Maps Is Hard

In general, finding feature maps for a corresponding kernel is hard.

**Example:** Radial Kernel,  $x \in \mathbb{R}$

► Consider



$$\phi(x) = (1, \frac{xe^{-x^2/2}}{\sqrt{1!}}, \frac{x^2e^{-x^2/2}}{\sqrt{2!}}, \dots, \frac{x^ne^{-x^2/2}}{\sqrt{n!}}, \dots)$$

► Then

$$\begin{aligned}\langle \phi(x_i), \phi(x_j) \rangle &= \sum_{n=0}^{\infty} \frac{x_i^n e^{-x_i^2/2}}{\sqrt{n!}} \frac{x_j^n e^{-x_j^2/2}}{\sqrt{n!}} \\ &= e^{-x_i^2/2 - x_j^2/2} \sum_{n=0}^{\infty} \frac{x_i^n x_j^n}{n!} = e^{-\|x_i - x_j\|^2/2} = K(x_i, x_j)\end{aligned}$$

Fortunately, we can use Kernels without knowing the feature maps.  
We'll talk about this after the midterm.

# Example



- ▶ Left: An SVM with a polynomial kernel of degree 3
- ▶ Right: An SVM with a radial kernel
- ▶ Either kernel is capable of capturing the decision boundary

# SVM with More than Two Classes

## One-Versus-One classification

- ▶ For  $K > 2$  classes, consider  $\binom{K}{2}$  pairs
- ▶ For example, we might compare the  $k$ th class, coded as  $+1$ , to the  $k'$ th class, coded as  $-1$
- ▶ We tally the number of times that the test observation is assigned to each of the  $K$  classes
- ▶ The final classification is performed by assigning the test observation to the class to which it was most frequently assigned in these  $\binom{K}{2}$  pairwise classifications

# SVM with More than Two Classes

## One-Versus-All classification

- ▶ We fit  $K$  SVMs, each time comparing one of the  $K$  classes to the remaining  $K - 1$  classes.
- ▶ Let  $\beta_{0k}, \beta_{1k}, \dots, \beta_{pk}$  denote the parameters that result from fitting an SVM comparing the  $k$ th class (coded as +1) to the others (coded as -1).
- ▶ Let  $x_0$  denote a test observation
- ▶ We assign  $x^*$  to the class with maximum

$$\beta_{0k} + \beta_{1k}x_1^* + \dots + \beta_{pk}x_p^*$$



# Khan - Gene Expression Data

- ▶ Gene expression measurements for four cancer types of small round blue cell tumors.
- ▶ For each tissue sample, 2308 gene expression measurements are available.

```
> library(ISLR)
> names(Khan)
[1] "xtrain" "xtest"  "ytrain" "ytest"
> dim(Khan$xtrain)
[1] 63 2308
> dim(Khan$xtest)
[1] 20 2308
> length(Khan$ytrain)
[1] 63
> length(Khan$ytest)
[1] 20
```

# Khan - Gene Expression Data

- ▶ The training and test sets consist of 63 and 20 observations respectively
- ▶ Belong to 4 cancer types

```
> table(Khan$ytrain)
 1  2  3  4
 8 23 12 20
> table(Khan$ytest)
 1 2 3 4
 3 6 6 5
```

# Khan - SVM Classification

Perfect fit with linear kernels (!) - no training errors.

Is this a surprise?

```
> dat=data.frame(x=Khan$xtrain, y=as.factor(Khan$ytrain))
> out=svm(y~., data=dat, kernel="linear",cost=10)
> summary(out)
Call:
svm(formula = y ~ ., data = dat, kernel = "linear",
     cost = 10)
Parameters:
  SVM-Type:  C-classification
 SVM-Kernel: linear
       cost: 10
    gamma: 0.000433
Number of Support Vectors: 58
( 20 20 11 7 )
Number of Classes: 4
Levels:
 1 2 3 4
> table(out$fitted, dat$y)
```

	1	2	3	4
1	8	0	0	0
2	0	23	0	0
3	0	0	12	0
4	0	0	0	20

# Khan - SVM Test Performance

## Two testing errors

```
> dat.te=data.frame(x=Khan$xtest, y=as.factor(Khan$ytest))  
> pred.te=predict(out, newdata=dat.te)  
> table(pred.te, dat.te$y)
```

```
pred.te 1 2 3 4  
      1 3 0 0 0  
      2 0 6 2 0  
      3 0 0 4 0  
      4 0 0 0 5
```

## Reading on Support Vector Machines

**ISL:** Chapter 9. In particular, read Section 9.6 on experiments in R, including: e1071 library, "Khan" gene expression data (and ROC curves)

**ESL:** Chapter 12

**Paper:** Vladimir N. Vapnik (1999), "An Overview of Statistical Learning"

**Homework:** Start working on the final project.

**Optional reading: Optimization** - book

Convex Optimization, Stephen Boyd and Lieven Vandenberghe,  
Cambridge University Press, 2004

Free download:

[http://stanford.edu/~boyd/cvxbook/bv\\_cvxbook.pdf](http://stanford.edu/~boyd/cvxbook/bv_cvxbook.pdf)

Today's presentation can be found in Chapter 5.