

EECS E6690: SL for Bio & Info

Lecture 11: Dimensionality Reduction, Information Ranking, Association Analysis, and Social Networks

Prof. Predrag R. Jelenković
Time: Tuesday 4:10-6:40pm
1127 Seeley W. Mudd Building

Dept. of Electrical Engineering
Columbia University , NY 10027, USA
Office: 812 Schapiro Research Bldg.
Phone: (212) 854-8174
Email: predrag@ee.columbia.edu
URL: <http://www.ee.columbia.edu/~predrag>

Unsupervised Learning

There is no input-output relationship, $y = f(x)$, i.e. **there is no y** .

- ▶ We only have a bunch of points (x_1, x_2, \dots, x_n)
- ▶ The problem has less structure
- ▶ We are trying to discover a structure - maybe more interesting

Typical questions and approaches

- ▶ Clustering - grouping data points
- ▶ Principal Component Analysis (PCA): used of preprocessing and visualization
- ▶ **Ranking: e.g., Google's PageRank algorithm**
- ▶ **Association Rules - Market Basket Analysis**
discovering relationships between data points
- ▶ Community Detection or Graph Clustering: e.g., discovering communities on Facebook

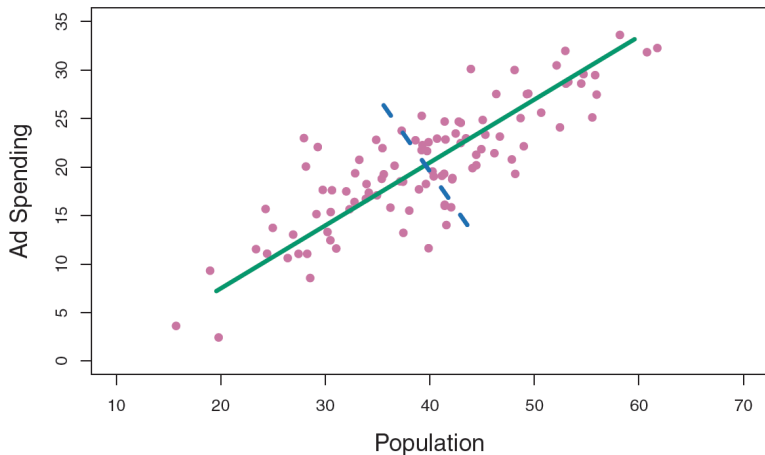
Last lecture: Principal Components Analysis (PCA)

Principal Components Analysis - unsupervised approach (no use of Y)

- ▶ Finding directions along which data is located - highly variable
- ▶ These directions define lines/subspaces that are close to the "data cloud"
- ▶ Can be used for visualization/understanding data
- ▶ Can be used as preprocessing for supervised learning

PCA: Example

Two PC-s: Solid line: First PC; Dashed line: Second PC



PCA: Good for High Dimensional Data - Large p

- ▶ Dimensionality reduction - e.g. gene expression data
- ▶ Assume we compute M principal components
- ▶ The best M -dimensional approximation to x_{ij}

$$x_{ij} = \sum_{m=1}^M z_{im} \phi_{jm}$$

How good is PC approximation?

- ▶ Proportion of Variance Explained (PVE)
- ▶ For each component, m , PVE is equal to

$$\frac{\sum_{i=1}^n \left(\sum_{j=1}^p \phi_{jm} x_{ij} \right)^2}{\sum_{ij} x_{ij}^2}$$

PCA is an Eigen-Decomposition Problem

- ▶ In general finding k -principal components is equivalent to finding a $k \times p$ matrix such that $\mathbf{z}_i = W\mathbf{x}_i$

- ▶ We also want linear recovery, i.e., to find a matrix U , such that $\hat{\mathbf{x}}_i = U\mathbf{z}_i = UW\mathbf{x}_i \approx \mathbf{x}_i$

- ▶ Hence, this reduces to an optimization problem

$$\arg \min_{W,U} \sum_{i=1}^n \|\mathbf{x}_i - UW\mathbf{x}_i\|^2 \quad (1)$$

Theorem Let $A = \sum_{i=1}^n \mathbf{x}_i \mathbf{x}_i^\top$ and let $\mathbf{u}_1, \dots, \mathbf{u}_k$ be the k leading eigenvectors of A . Then, the solution to the PCA problem is to set the columns of U to be $\mathbf{u}_1, \dots, \mathbf{u}_k$ and to set $W = U^\top$.

Efficient Solution for High-Dimensional Data: $p \gg n$

The complexity of the previously described PCA solution is $O(p^3)$:

- ▶ $O(np^2)$ is the complexity of computing the matrix $A = \sum_{i=1}^n \mathbf{x}_i \mathbf{x}_i^\top$.
- ▶ $O(p^3)$ is the complexity of eigendecomposition of A .

Now, consider an alternative calculation:

- ▶ Rewrite $A = X^\top X$, where X is $n \times p$ matrix with i th row \mathbf{x}_i^\top
- ▶ Consider $K = XX^\top$, $n \times n$, dot-product matrix (linear kernel), whose (i, j) element is $\langle \mathbf{x}_i, \mathbf{x}_j \rangle$
- ▶ Suppose \mathbf{u} is an eigenvector of K :

$$K\mathbf{u} = \lambda\mathbf{u}$$

- ▶ Multiplying the preceding equality by X^\top : $X^\top XX^\top \mathbf{u} = \lambda X^\top \mathbf{u}$
- ▶ Implying: $A(X^\top \mathbf{u}) = \lambda(X^\top \mathbf{u})$
- ▶ Hence, $\frac{X^\top \mathbf{u}}{\|X^\top \mathbf{u}\|}$ is the eigenvector of A
- ▶ We can thus calculate the PCA solution from eigendecomposition of K with complexity $O(n^3 + n^2p)$

Kernel PCA

- ▶ Expand basis $x \rightarrow \phi(x)$, where $\phi(x)$ is in some Hilbert space
- ▶ Now, PCA is based on **orthogonal projections**

$$\langle \phi(x_i), \phi(x) \rangle$$

- ▶ Which, based on the preceding slide, can be done by performing eigen decomposition of the kernel matrix $K = \{\langle \phi(x_i), \phi(x_j) \rangle\}$
- ▶ Moreover, recall that our approximation functions are of the form

$$g_1(x) = \sum_{i=1}^n \phi_{i1} k(x, x_i)$$

- ▶ Hence, we can think of finding the first principle component of Kernel PCA as the variance maximization

$$\max_{g_1 \in \mathcal{H}_k} \text{Sample Var}(g_1(X)) \quad \text{subject to} \quad \|g_1\|_{\mathcal{H}_k} = 1$$

- ▶ Now the approximation "directions" are not lines anymore

Example: Dimensionality Reduction of Images



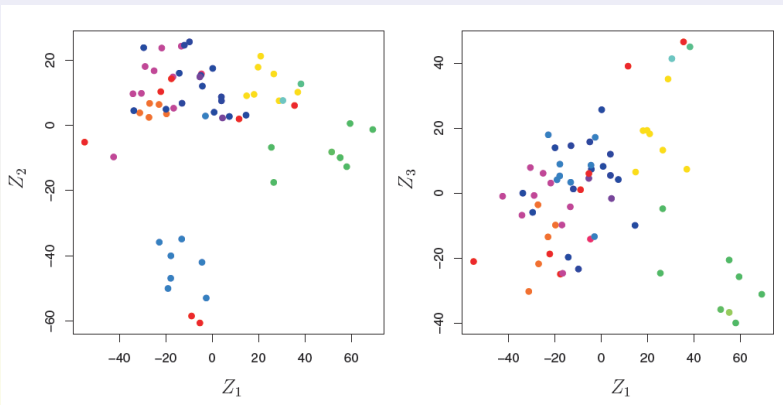
- ▶ Top left: 50×50 pixel image: $p = 2,500$ dimensions
- ▶ Top right: Reconstruction with $k = 10$ principal components:
 $10 \ll 2,500(!)$

PCA on the NCI60 Data

- ▶ First *scale* the data
 `> pr.out=prcomp(nci.data, scale=TRUE)`
- ▶ Then, **assign a color** to each of the 64 cancer cell lines
 `Cols=function(vec){
 + cols=rainbow(length(unique(vec)))
 + return(cols[as.numeric(as.factor(vec))]) }`
- ▶ We now can plot the principal component score vectors
 `> par(mfrow=c(1,2))
 > plot(pr.out$x[,1:2], col=Cols(nci.labs), pch=19,
 xlab="Z1",ylab="Z2")
 > plot(pr.out$x[,c(1,3)], col=Cols(nci.labs), pch=19,
 xlab="Z1",ylab="Z3")`

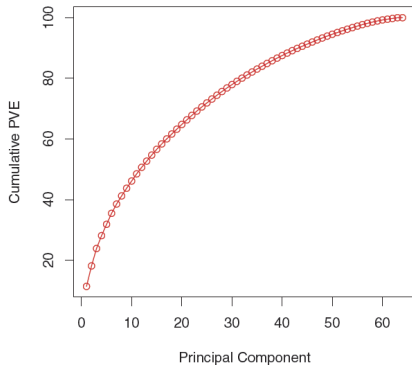
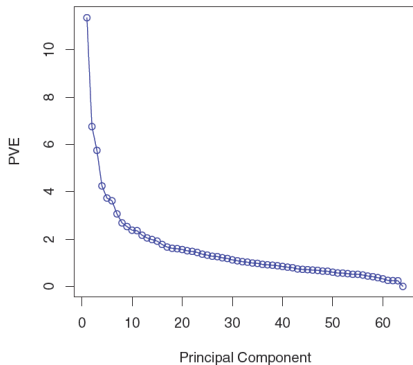
PCA on the NCI60 Data

- ▶ Observations belonging to a single cancer type tend to lie near each other in this low-dimensional space.
- ▶ Not possible to visualize data without PCA, $\binom{6832}{2}$ scatter plots (!)



Principal Value Explained (PVE) on NCI60 Data

- ▶ First 7 PC-s explain 40% of data
- ▶ 70 PC-s explain 100% of data
- ▶ Compare 70 to $p = 6832(!)$



Other Dim-Reduction Techniques: Random Projections

- ▶ Let W be a random $m \times p, m < p$ matrix, then the transformation

$$\mathbf{x} \rightarrow W\mathbf{x}$$

is a **random projection**.

- ▶ Simple compression technique that works remarkably well.
- ▶ It is based on Johnson-Lindenstrauss lemma, which shows that random projections do not distort Euclidian distance a lot, i.e.,

$$\|\mathbf{x}_1 - \mathbf{x}_2\| \approx \|W\mathbf{x}_1 - W\mathbf{x}_2\|,$$

which is equivalent to showing that $\|W\mathbf{x}\|/\|\mathbf{x}\| \approx 1$.

Lemma Fix $\mathbf{x} \in \mathbb{R}^p$. Let W be $m \times p$ random matrix with independent standard normal components. Then, for every $\epsilon \in (0, 3)$,

$$\mathbb{P} \left[\left| \frac{\|W\mathbf{x}/\sqrt{m}\|^2}{\|\mathbf{x}\|^2} - 1 \right| > \epsilon \right] \leq 2e^{-\epsilon^2 m/6}.$$

Other Dim-Reduction Techniques: Random Projections

Proof: Without loss of generality assume $\|\mathbf{x}\| = 1$. Hence, it remains to prove

$$\mathbb{P}[(1 - \epsilon)m \leq \|W\mathbf{x}\|^2 \leq (1 + \epsilon)m] \geq 1 - 2e^{-\epsilon^2 m/6}$$

To this end, if \mathbf{w}_i is the i th row of W , then $\langle \mathbf{w}_i, \mathbf{x} \rangle$ is normal with mean zero and

$$\sigma^2 = \sum_{j=1}^p x_j^2 = \|\mathbf{x}\|^2 = 1$$

Hence, $\|W\mathbf{x}\|^2$ has χ_m^2 distribution since

$$\|W\mathbf{x}\|^2 = \sum_{i=1}^m \langle \mathbf{w}_i, \mathbf{x} \rangle^2,$$

and $\langle \mathbf{w}_i, \mathbf{x} \rangle$ are independent standard normal variables. Thus, the claim of the lemma is equivalent to

$$\mathbb{P}[(1 - \epsilon)m \leq \chi_m^2 \leq (1 + \epsilon)m] \geq 1 - 2e^{-\epsilon^2 m/6},$$

which can be derived using Chernoff's bound, e.g., see Lemma B.12 in Shai Shalev-Shwartz and Shai Ben-David book.

Other Dim-Reduction Techniques: Compressed Sensing

Compressed Sensing:

- ▶ Also a linear technique

$$\mathbf{x} \rightarrow W\mathbf{x}$$

- ▶ Matrix W has a special property: Reconstructed Isoperimetric Property (RIP): A matrix W is (ϵ, s) RIP if for all $\|\mathbf{x}\| \neq 0$, such that $\|\mathbf{x}\|_0 \leq s$,

$$\left| \frac{\|W\mathbf{x}\|^2}{\|\mathbf{x}\|^2} - 1 \right| \leq \epsilon$$

- ▶ A random (!) $m \times p$ matrix is likely to satisfy the RIP condition provided that $m > s \log(p)$.
- ▶ Reconstruction solves a linear program, and it can be computed in polynomial time.

See Section 23.3 of Shai Shalev-Shwartz and Shai Ben-David book for more details.

Information Retrieval from Large Related Data: the Web

- ▶ Suppose we have a large amount of related data - e.g. the Web
- ▶ Relationship between data items, e.g. Web pages, can be expressed as a graph
- ▶ Information retrieval - Search: suppose we are looking for all data items (Web pages) that contain some information/keywords
- ▶ Typically the search will return a large number of results
- ▶ Hence, we **need to rank them**
- ▶ How?
- ▶ Example: The Google PageRank Algorithm

The Google PageRank Algorithm

- ▶ Idea: A page should be important if
 - ▶ It is pointed to by highly ranked pages
 - ▶ It is pointed/connected to by a lot of pages

Definition

- ▶ N - total number of pages
- ▶ Graph connectivity - incidence matrix
 $L_{ij} = 1$ if page j points to page i
- ▶ $c_j = \sum_{i=1}^N L_{ij}$ - number of pages that point to by page j - out-degree of j
- ▶ Then, PageRank assigns a rank, R_i , to page i according to

$$R_i = (1 - d) + d \sum_{j=1} \frac{L_{ij}}{c_j} R_j,$$

where $0 < d < 1$ is a constant set to $d = 0.85$

PageRank Algorithm: Explanation

Let us look back at the equation

$$R_i = (1 - d) + d \sum_{j=1}^N \frac{L_{ij}}{c_j} R_j,$$

where $0 < d < 1$ is a constant set to $d = 0.85$

- ▶ d - ensures that each page gets at least rank $1 - d$
- ▶ d - also ensures that the preceding system of linear equations has a unique solution
- ▶ If some of the j -s that point to i have high rank, R_j , then R_i will be high, i.e.
 - ▶ Page i is important if some of the j -s are important
- ▶ $1/c_j$ is designed to prevent spamming, i.e., if j liberally points to a lot of pages, then its contribution to i 's ranking should count less
- ▶ If i has a lot of neighbors, i.e., large $\sum_{j=1}^N L_{ij}$, then its rank is large

PageRank: Solution

Let us introduce matrix notation

- ▶ e - vector of N ones
- ▶ D_c - diagonal matrix with elements c_j
- ▶ \mathbf{R} - column vector of page ranks (R_1, R_2, \dots, R_N) .
- ▶ Then, PageRank equations can be written as

$$\mathbf{R} = (1 - d)e + d\mathbf{L}D_c^{-1}\mathbf{R}$$

- ▶ Introducing the normalization that the average page rank is one

$$e^\top \mathbf{R} = N$$

we can write the preceding linear system as

$$\mathbf{R} = \left[(1 - d) \frac{ee^\top}{N} + d\mathbf{L}D_c^{-1} \right] \mathbf{R} = \mathbf{A}\mathbf{R}$$

where matrix \mathbf{A} is the expression in square brackets

PageRank: Solution

Connection to Markov chains

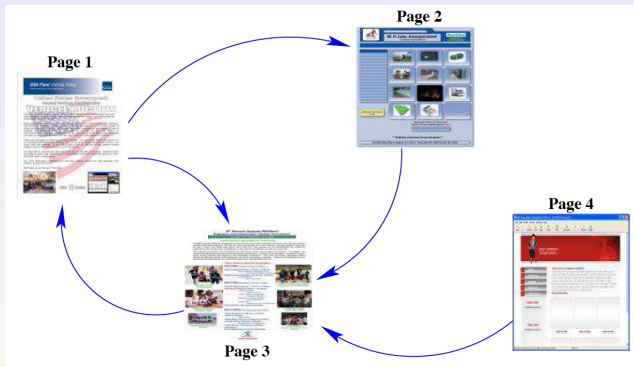
- ▶ \mathbf{A} - is a stochastic matrix
 - ▶ Largest eigenvalue of \mathbf{A} is 1
- ▶ \mathbf{R} can be computed iteratively by power method, starting with $\mathbf{R} = \mathbf{R}_0$, and performing

$$\mathbf{R}_k \leftarrow \mathbf{A}\mathbf{R}_{k-1}; \quad \mathbf{R}_k \leftarrow N \frac{\mathbf{R}_k}{\mathbf{e}^\top \mathbf{R}_k}$$

Random surfer interpretation

- ▶ The surfer performs a random walk on the Web choosing among the outgoing links at random
- ▶ The surfer can also jump with probability $(1 - d)$ to a random page on the Web
- ▶ Then, R_i represents the probability that a surfer is found at page i .

PageRank: Example



$$\mathbf{L} = \begin{bmatrix} 0 & 0 & 1 & 0 \\ 1 & 0 & 0 & 0 \\ 1 & 1 & 0 & 1 \\ 0 & 0 & 0 & 0 \end{bmatrix} \quad \mathbf{c} = [2 \quad 1 \quad 1 \quad 1] \quad \mathbf{R} = [1.49 \quad 0.78 \quad 1.58 \quad 0.15]$$

PageRank: Stochastic Analysis

- ▶ Numerical solution for R doesn't reveal qualitative behavior
- ▶ Measurements showed that R follows power law behavior $\sim 1/i^\alpha$
- ▶ Why?

In Jelenković & Olvera-Cravioto (2010), we introduced a stochastic formulation

$$R = Q + \sum_{i \in \mathcal{N}} C_i R_i$$

where

- ▶ R is a typical, random, page on the Web
- ▶ \mathcal{N} - typical, random, number of neighbors that are pointing to this page
- ▶ R_i - typical rank of these neighbors
- ▶ C_i - random variables taking into account d/c_j
- ▶ Q - random variable taking into account d

PageRank: Stochastic Solution

Main implications of Jelenković & Olvera-Cravioto (2010) and later work

- ▶ Explain why R follows power law behavior $\sim 1/i^\alpha$
- ▶ Known fact - Web graph is a power law graph, i.e. $|\mathcal{N}|$ - has a power law distribution
- ▶ R can either be large because
 - ▶ it is pointed to by a large number of lower-ranked pages
 - ▶ or, it is pointed to by a small number of lower-ranked pages
- ▶ Reveals a weakness: the factor $1/c_j$ doesn't prevent spamming

Improvements

- ▶ Prevent spamming
- ▶ Design other ranking schemes, e.g. personalized Web search

Association Rules: Market Basket Analysis

- ▶ Popular technique for mining commercial data bases
 - ▶ Amazon or Walmart customer data bases
- ▶ $X = (X_1, \dots, X_p)$
- ▶ p - total number of items in the store - **very large**
- ▶ Often $X_i \in \{0, 1\}$ - so $X = (X_1, \dots, X_p)$ represents the times that a customer bought
 - ▶ $\mathbf{x}_i = (x_{i1}, \dots, x_{ip})$ represents items that customer i bought
"market basket"
- ▶ Those variables that frequently have joint values of one represent items that are frequently purchased together.
- ▶ General problem - find vectors v_i , such that $\mathbb{P}[X = v_i]$ is large
- ▶ This problem is impossibly difficult
- ▶ Can we compute the histogram of X ?
 p large - never enough data

Association Rules: Simplification

- ▶ Instead of finding **values of v** for which $\mathbb{P}[X = v]$ is large
 - ▶ **Find regions** of the X -space with high probability relative to its size or support
 - ▶ Let s_j be a **subset of values** that X_j can take, then we try to find these subsets s_j for which

$$\mathbb{P}[\cap_{j=1}^p \{X_j \in s_j\}]$$

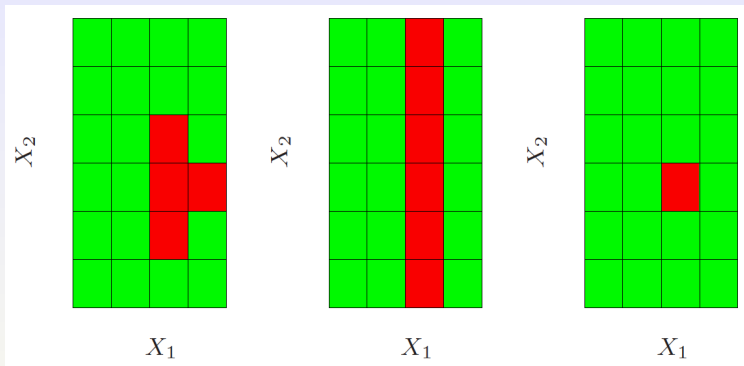
- ▶ Even this is not feasible for $p \approx 10^4$ and $N \approx 10^8$
- ▶ Note that if s_j = entire set of values for X_j , then feature j is excluded
- ▶ Hence, we can only focus on a subset of items $\mathcal{J} \subset \{1, \dots, p\}$ and try to maximize

$$\mathbb{P}[\cap_{j \in \mathcal{J}} \{X_j \in s_j\}]$$

or even simpler take each s_j to be a single item

$$\mathbb{P}[\cap_{j \in \mathcal{J}} 1_{\{X_j = v_j\}}]$$

Association Rules: Example



- ▶ The red squares indicate areas of high density.
- ▶ To simplify, we assume that the derived subset corresponds to either a single value of an input or all values.
- ▶ With this assumption we could find either the middle or right pattern, but not the left one.

Turn Everything to Binary Variables

- ▶ Introducing dummy variables can turn X into binary vectors Z
- ▶ Recall \mathcal{S}_j - all values that X_j can take
- ▶ Let $K = \sum_{j=1}^p |\mathcal{S}_j|$
- ▶ Then $Z_k = 1$ is item k , $1 \leq k \leq K$ is present in the basket
- ▶ $\mathcal{K} \subset \{1, \dots, K\}$ - item set, then we maximize

$$\mathbb{P}[\prod_{k \in \mathcal{K}} Z_k = 1]$$

- ▶ The preceding probability can be estimated as

$$\mathbb{P}[\prod_{k \in \mathcal{K}} Z_k = 1] \approx \frac{1}{N} \sum_{i=1}^N \prod_{k \in \mathcal{K}} z_{ik}$$

this is called the "**support**" or "**prevalence**", $T(\mathcal{K})$, of the item set \mathcal{K}
 z_{ik} is the value of Z_k for the i -th case/customer

- ▶ For a threshold t , find all item sets \mathcal{K}_l , for which $T(\mathcal{K}_l) > t$

The Apriori Algorithm

- ▶ Agrawal et al. (1995) - exploits
 - ▶ Choose a threshold t such that sets \mathcal{K}_l , for which $T(\mathcal{K}_l) > t$, have small number of items
 - ▶ If $\mathcal{L} \in \mathcal{K}$, then $T(\mathcal{L}) \geq T(\mathcal{K})$
- ▶ First pass: Find all single items whose support is bigger than the threshold (other items are discarded)
- ▶ Second pass: computes the support of all item sets of size two that can be formed from pairs of the single items surviving the first pass.
- ▶ Continue this procedure for sets of $|\mathcal{K}| = m, m = 3, 4, \dots$ items, by only considering only those from the previous, $m - 1$, pass with those retained from the first pass.
- ▶ Passes over the data continue until all candidate rules from the previous pass have support less than the specified threshold.
- ▶ The Apriori algorithm requires only one pass over the data for each value of $|\mathcal{K}|$
- ▶ If the data are sufficiently sparse (or t is high enough), then the process will terminate in reasonable time even for huge data sets.

The Apriori Algorithm

- ▶ Each high support item set \mathcal{K} returned by the Apriori algorithm is cast into a set of association rules.
- ▶ The items $Z_k, k \in \mathcal{K}$ are partitioned into two disjoint subsets, $\mathcal{A} \cup \mathcal{B} = \mathcal{K}$, and written

$$\mathcal{A} \Rightarrow \mathcal{B}$$

- ▶ The first item subset \mathcal{A} is called the **antecedent**
- ▶ The second \mathcal{B} the **consequent**
- ▶ The **support** of the rule $T(\mathcal{A} \Rightarrow \mathcal{B}) =$ fraction of observations in the union of the antecedent and consequent, which is just the support of the item set \mathcal{K}
- ▶ $T(\mathcal{A} \Rightarrow \mathcal{B}) = T(\mathcal{K})$ is an estimate of the probability of simultaneously observing both item sets $\mathbb{P}(\mathcal{A} \cap \mathcal{B})$ in a randomly selected market basket.

The Apriori Algorithm

- ▶ The confidence or predictability

$$C(\mathcal{A} \Rightarrow \mathcal{B}) = \frac{T(\mathcal{A} \Rightarrow \mathcal{B})}{T(\mathcal{A})}$$

is an estimate of the conditional probability $\mathbb{P}(\mathcal{B}|\mathcal{A})$

- ▶ The "lift"

$$L(\mathcal{A} \Rightarrow \mathcal{B}) = \frac{C(\mathcal{A} \Rightarrow \mathcal{B})}{T(\mathcal{B})}$$

is an estimate of

$$\frac{\mathbb{P}(\mathcal{A} \cap \mathcal{B})}{\mathbb{P}(\mathcal{A})\mathbb{P}(\mathcal{B})}$$

The Apriori Algorithm: Example

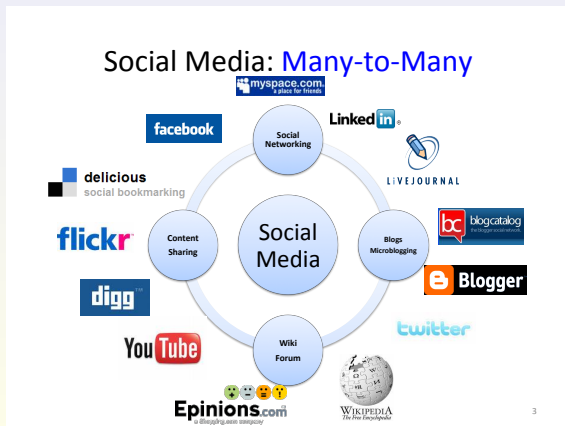
- ▶ $\mathcal{K} = \{\text{peanut butter, jelly, bread}\}$
- ▶ consider the rule $\{\text{peanut butter, jelly}\} \Rightarrow \{\text{bread}\}$.
- ▶ support value of 0.03 for this rule means that peanut butter, jelly, and bread appeared together in 3% of the market baskets.
- ▶ confidence of 0.82 for this rule implies that when peanut butter and jelly were purchased, 82% of the time bread was also purchased.
- ▶ If bread appeared in 43% of all market baskets then the rule $\{\text{peanut butter, jelly}\} \Rightarrow \{\text{bread}\}$ would have a lift of 1.95.

Goal: to find rules with **high support and confidence**

Social Networks

This part of the lecture is based on the book:
Community Detection and Mining in Social Media, L. Tang & H. Liu,
2010

- World Wide Web, Facebook, Twitter, Youtube, etc.



Characteristics of Social Media

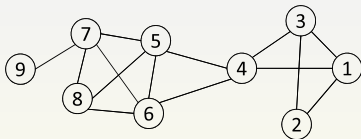
- ▶ Consumers become Producers
- ▶ Rich User Interaction
- ▶ User-Generated Contents
- ▶ Collaborative environment
- ▶ Collective Wisdom
- ▶ Long Tail



Social Media: Graph Representation

- ▶ Social Network: made of nodes (individuals or organizations) and edges that connect nodes in various relationships like friendship, kinship etc.

- Graph Representation



- Matrix Representation

Node	1	2	3	4	5	6	7	8	9
1	-	1	1	1	0	0	0	0	0
2	1	-	1	0	0	0	0	0	0
3	1	1	-	1	0	0	0	0	0
4	1	0	1	-	1	1	0	0	0
5	0	0	0	1	-	1	1	1	0
6	0	0	0	1	1	-	1	1	0
7	0	0	0	0	1	1	-	1	1
8	0	0	0	0	1	1	1	-	0
9	0	0	0	0	0	0	1	0	-

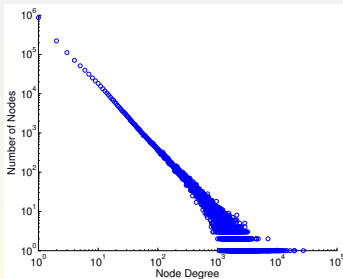
Graphs: Basic Concepts

- ▶ A : the adjacency matrix
- ▶ V : the set of nodes
- ▶ E : the set of edges
- ▶ v_i : a node v_i
- ▶ $e(v_i, v_j)$: an edge between node v_i and v_j
- ▶ N_i : the neighborhood of node v_i
- ▶ d_i : the degree of node v_i
- ▶ Geodesic distance: a shortest path between two nodes

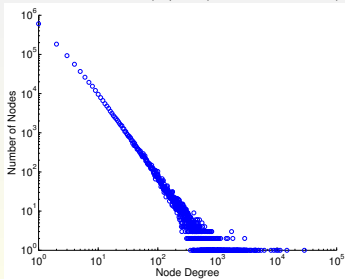
Properties of Large-Scale Social Networks

- ▶ Networks in social media are typically huge, involving millions of actors and connections.
- ▶ Large-scale networks in real world demonstrate similar patterns
 - ▶ Strong Community Structure
 - ▶ Small-world effect
 - ▶ Scale-free distributions: power laws - $1/i^\alpha$

Log-Log Plot of Power Law Distributions: $\log(1/i^\alpha) = -\alpha \log(i)$



Friendship Network in Flickr



Friendship Network in YouTube ¹³

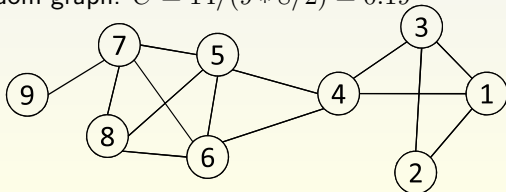
Measuring Community Structure

Clustering coefficient:

$$C_i = \frac{\# \text{ of closed triplets}}{\# \text{ of connected closed triplets}}$$

Example:

- ▶ $d_6 = 4, N_6 = \{4, 5, 7, 8\}$
 - ▶ $k_6 = 4 : e(4, 5), e(5, 7), e(5, 8), e(7, 8)$
 - ▶ $C_6 = 4 / (4 * 3 / 2) = 2 / 3$
 - ▶ Average clustering coefficient $C = (C_1 + C_2 + \dots + C_n) / n$
 - ▶ $C = 0.61$ for the example network
- In a random graph: $C = 14 / (9 * 8 / 2) = 0.19$



Some Social Network Challenges

Network Modeling

- ▶ Small-world effect (e.g., 6 degrees of separation)
- ▶ Power-law distribution (a.k.a. scale-free distribution)
- ▶ Community structure (high clustering coefficient)
- ▶ Common model: Preferential attachment
- ▶ Take a look at:
Directed random graphs with given degree distributions, N. Chen & M. Olvera-Cravioto, Stochastic Systems, 3 (1), 147-186, 2013.

Community Detection - Graph Clustering

Community Detection - Graph Clustering

A community is a set of nodes between which the interactions are (relatively) frequent

- ▶ A.k.a., group, cluster, cohesive subgroups, modules



Applications: Recommendation based communities, Network Compression, Visualization of a huge network

Clustering Based on Vertex Similarity

This is familiar(!)

Define node similarity/distance:

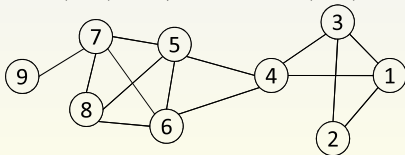
- ▶ Jaccard Similarity

$$d(v_i, v_j) = \text{Jaccard}(v_i, v_j) = \frac{|N_i \cap N_j|}{|N_i \cup N_j|}$$

- ▶ Cosine similarity

$$d(v_i, v_j) = \text{Cosine}(v_i, v_j) = \frac{\sum_k A_{ik} A_{jk}}{\sqrt{\sum_k A_{ik}^2 \sum_k A_{jk}^2}}$$

$$\text{Jaccard}(4, 6) = 1/7, \quad \text{Cosine}(4, 6) = 1/4$$



Then, we can run any known clustering algorithm, e.g. K-means

Recall K-Means

What are centroids on a graph? - define properly

- ▶ **Centroid of a graph:** A vertex (or set vertices) u where the greatest distance $d(u, v)$ to other vertices v is minimal.

K-Means Algorithm

1. Initialization: Randomly assign a number, from 1 to K , to each of the nodes. (Or, select K nodes randomly to be centroids.)
2. Iterate until the cluster assignments stop changing:
 - (a) For each of the K clusters, compute the cluster centroid.
 - (b) Assign each observation to the cluster whose centroid is closest.

Divisive Hierarchical Clustering: Bottom up

Divisive clustering

- ▶ Partition nodes into several sets
- ▶ Each set is further divided into smaller ones

Example: recursively remove the weakest tie

- ▶ Find the edge with the least strength
- ▶ Remove the edge and update the corresponding strength of each edge

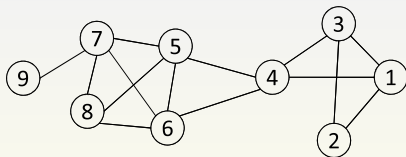
Recursively apply the above two steps until a network is decomposed into desired number of connected components.

Each component forms a community.

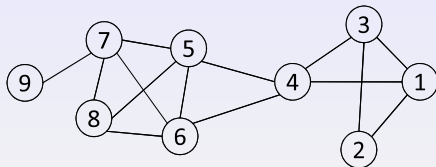
Strength of a Tie: Edge Betweenness

Edge betweenness: the # of shortest paths that pass along with the edge

The edge betweenness of $e(1, 2)$ is 4, as:
all the shortest paths from 2 to $\{4, 5, 6, 7, 8, 9\}$ have to either pass $e(1, 2)$ or $e(2, 3)$, and $e(1, 2)$ is the shortest path between 1 and 2



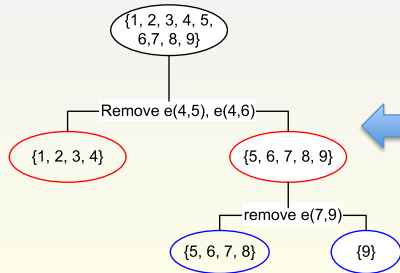
Example



Initial betweenness value

Table 3.3: Edge Betweenness

	1	2	3	4	5	6	7	8	9
1	0	4	1	9	0	0	0	0	0
2	4	0	4	0	0	0	0	0	0
3	1	4	0	9	0	0	0	0	0
4	9	0	9	0	10	10	0	0	0
5	0	0	0	10	0	1	6	3	0
6	0	0	0	10	1	0	6	3	0
7	0	0	0	0	6	6	0	2	8
8	0	0	0	0	3	3	2	0	0
9	0	0	0	0	0	0	8	0	0

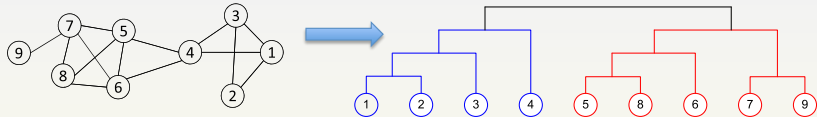


After remove $e(4,5)$, the betweenness of $e(4, 6)$ becomes 20, which is the highest;

After remove $e(4,6)$, the edge $e(7,9)$ has the highest betweenness value 4, and should be removed.

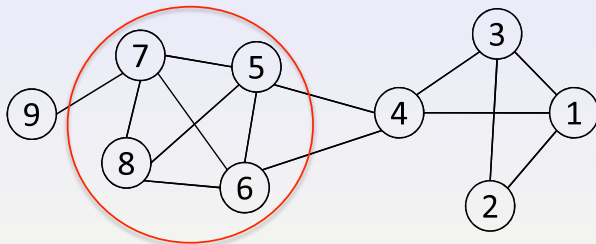
Agglomerative Hierarchical Clustering

- ▶ Initialize each node as a community
- ▶ Merge communities successively into larger communities following a certain criterion



Other Ideas: Using Cliques

Clique: a maximum complete subgraph in which all nodes are adjacent



- ▶ NP-hard to find the maximum clique in a network
- ▶ Straightforward implementation to find cliques is very expensive in time complexity
- ▶ There are heuristics: Find a clique of size k , and prune all nodes with $< k - 1$ neighbors
- ▶ **Idea for clustering:** Use cliques as seeds to form larger communities

Approximating Cliques: Density Based Groups

- ▶ The group-centric criterion requires the whole group to satisfy a certain condition
 - ▶ E.g., the group density \geq a given threshold
- ▶ A subgraph $G_s(V_s, E_s)$ is a γ -dense quasi-clique if

$$\frac{|E_s|}{|V_s|(|V_s| - 1)/2} \geq \gamma$$

- ▶ A similar strategy to that of cliques can be used
 - ▶ Find a maximal quasi-clique, say, of size k
 - ▶ Remove nodes with degree $< \gamma k$

Dimensionality Reduction on Graphs

- ▶ Map nodes into a low-dimensional space such that:
 - ▶ The proximity between nodes is preserved in the new space
 - ▶ Solutions usually involve eigen-decomposition
 - ▶ Then perform clustering, e.g. k-means

Reading on Community Detection and Social Networks

- ▶ Book: Community Detection and Mining in Social Media, Lei Tang & Huan Liu, 2010
- ▶ Free download: <https://www.morganclaypool.com/doi/abs/10.2200/S00298ED1V01Y201009DMK003>
- ▶ Other book resources: <http://dmml.asu.edu/cdm/>

Comprehensive R-Package for Graph Analysis

- ▶ iGraph - <http://igraph.org/r/#docs>
- ▶ R manual: <http://igraph.org/r/doc/igraph.pdf>
- ▶ Various tutorials for iGraph in R available online

Reading on PageRank

ESL: Section 14.10

Paper P.R. Jelenković and M. Olvera-Cravioto, Information ranking and power laws on trees, *Advances in Applied Probability*, 42 (4), pp. 1057-1093, 2010.

Reading on Association Rules

ESL: Section 14.2

Paper Agrawal et al., Fast discovery of association rules, *Advances in Knowledge Discovery and Data Mining*, AAAI/MIT Press, Cambridge, MA, 1995.

Dimensionality Reduction: Chapter 23 in:

- ▶ Shai Shalev-Shwartz and Shai Ben-David, *Understanding Machine Learning: From Theory to Algorithms*, 2014. <https://www.cs.huji.ac.il/~shais/UnderstandingMachineLearning/understanding-machine-learning-theory-algorithms.pdf>

Final Project

- ▶ Groups of 4 students
- ▶ **Deliverables:**
 - ▶ **Paper** - 15 - 20 pages
 - ▶ **10min Presentation** with about 10 slides
- ▶ **Due: likely Dec 20, 11:59pm**
- ▶ One slot for presentations on **Tue, Dec 13, 4:10-6:40pm**
Other slots to be scheduled on **Friday, Dec 16**
- ▶ **Suggested Data Repositories:**
 - ▶ UC Irvine Machine Learning Repository
<https://archive.ics.uci.edu/ml/datasets.html>
 - ▶ Datasets supported by Bioconductor:
<http://www.bioconductor.org/packages/release/data/experiment/>
 - ▶ GEO (Gene Expression Omnibus) Data Repository
<https://www.ncbi.nlm.nih.gov/geo/>

Final Project

Final Paper Outline:

1. **Introduction**: e.g., describe the application area, problems considered, etc
2. **Data set(s) and paper(s)**: e.g., describe data in detail, what was done in the paper(s), common stat/machine learning tools, etc
3. **Reproduce the results from the paper(s)**
4. **Try different techniques learned in class, or propose new ones**
5. **Discussion and conclusion**: e.g., compare different techniques, pros and cons, future work, etc

General

- ▶ Document software well
- ▶ Use R, unless there is a compelling reason to use Python, e.g., you implemented a new algorithm
- ▶ **Academic Honesty** - do not plagiarize, as the papers will be submitted via **Turnitin**, which automatically checks for **plagiarization**.

Have fun and GOOD LUCK!