

EECS E6690: Statistical Learning for Biological and Information Systems Lecture 4: PCA, Nonlinear Models and Model Validation

Prof. Predrag R. Jelenković
Time: Tuesday 4:10-6:40pm

Dept. of Electrical Engineering
Columbia University , NY 10027, USA
Office: 812 Schapiro Research Bldg.
Phone: (212) 854-8174
Email: predrag@ee.columbia.edu
URL: <http://www.ee.columbia.edu/~predrag>

Last lecture: Linear Model Selection and Regularization

Motivation: Find the smallest/simplest model

- ▶ Problem with complex models
 - ▶ Overfitting (especially for $n \leq p$)
 - ▶ High testing error
- ▶ Model interpretability
 - ▶ Hard to interpret model with many predictors
 - ▶ Focus on most important variables
- ▶ Approaches for reduction:
 - ▶ Subset selection
 - ▶ Shrinkage methods
 - ▶ Dimension reduction techniques (today)

Last lecture: Bias-Variance trade-off

- ▶ “Test error”: test variable $y_0 = f(x_0) + \epsilon_0$

$$\begin{aligned}\text{Err}(x_0) &= \mathbb{E} \left(y_0 - \hat{f}(x_0) \right)^2 \\ &= \mathbb{E} \left(f(x_0) + \epsilon - \hat{f}(x_0) \right)^2 \\ &= \sigma^2 + \mathbb{E} \left(f(x_0) - \mathbb{E} \hat{f}(x_0) - \hat{f}(x_0) + \mathbb{E} \hat{f}(x_0) \right)^2 \\ &= \sigma^2 + \left(f(x_0) - \mathbb{E} \hat{f}(x_0) \right)^2 + \text{Var}(\hat{f}(x_0)) \\ &= \sigma^2 + \left(\text{Bias}(\hat{f}(x_0)) \right)^2 + \text{Var}(\hat{f}(x_0))\end{aligned}$$

- ▶ For linear models, **sample estimation** leads to
(see equation (7.12) in the [ESL] book)

$$\frac{1}{n} \sum_{i=1}^n \text{Err}(x_i) = \sigma^2 + \frac{1}{n} \sum_{i=1}^n \left(\text{Bias}(\hat{f}(x_i)) \right)^2 + \frac{p}{n} \sigma^2$$

- ▶ Increasing p reduces the bias but increases the variance

Last lecture: Subset selection


Deciding on the important variables:

- ▶ Need a criteria that balance training error and model size
- ▶ Several approaches
 - ▶ All (or best) subsets selection
 - ▶ Consider all 2^p models
 - ▶ Infeasible when p is large
 - ▶ Forward selection
 - ▶ Start with a null model – no predictors
 - ▶ Add predictors one-by-one
 - ▶ Stopping criterion
 - ▶ Backward selection
 - ▶ Start with a full model – p predictors
 - ▶ Eliminate predictors one-by-one
 - ▶ Stopping criterion

Last lecture: Best subset selection

```
# Making the plot
plot( 0:8, prostate.models.best.rss, ylim=c(0, 100),
      type="b", xlab="Subset size k", ylab="Residual Sum-of-Square", col="red2" )
points( prostate.models.size, prostate.models.rss, pch=20, col="dodgerblue",cex=0.7 )
```

Redo this with "glmnet" library: glmnet() function does both ridge ($\alpha=0$) and lasso ($\alpha=1$), and in-between.



Last lecture: Suboptimal (greedy) selection

Forward (or backward) stepwise selection

- ▶ Reduce computational complexity by forfeiting optimality
- ▶ Algorithm
 - ▶ Let \mathcal{M}_0 denote the null model (no predictors, sample mean prediction)
 - ▶ for $k = 0, 1, \dots, p-1$
 - ▶ Fit all $p-k$ models that augment the predictors in \mathcal{M}_k with one additional predictor
 - ▶ Let \mathcal{M}_{k+1} be the best of these $p-k$ models in terms of the smallest RSS (equivalently the largest R^2)
 - ▶ Select the best model from among $\mathcal{M}_0, \mathcal{M}_1, \dots, \mathcal{M}_p$ using some criterion
- ▶ Greedy
- ▶ Number of models is $1 + p(p+1)/2$
- ▶ If $p > n$, we can construct $\mathcal{M}_0, \dots, \mathcal{M}_n$ models only

Last lecture: Regularization (shrinkage) methods

- ▶ An alternative to subset selection
- ▶ Idea: Regularize/constrain coefficients
- ▶ Two widely-used methods:
 - ▶ Ridge regression
 - ▶ LASSO (least absolute shrinkage and selection operator)

Last lecture: Ridge

- ▶ OLS: minimize RSS

$$\text{RSS} = \sum_{i=1}^n \left(y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{i,j} \right)^2$$

- ▶ Ridge regression: minimize (RSS + shrinkage penalty)


$$\text{RSS} + \lambda \sum_{j=1}^p \beta_j^2$$

- ▶ λ is a tuning parameter: β_j^λ for each λ
- ▶ β_0 is not in the penalty

Prostate: Ridge regression

```
# Calling simple.ridge() function which is part of "ElemStatLearn" package
# Ridge functions in other packages:
# MASS: lm.ridge()
# mda : gen.ridge()
prostate.ridge <- simple.ridge( train[,1:8], train[,9], df=1:8 )

# plot
matplot( prostate.ridge$df, t(prostate.ridge$beta), type="b",
         col="blue", pch=20, ylab="coefficients" )
```



Last lecture: Lasso

- ▶ Ridge regression: Still p (shrunk) predictors
- ▶ Inference

- ▶ Lasso: minimize


$$\text{RSS} + \lambda \sum_{j=1}^p |\beta_j|$$

- ▶ Rationale for absolute values: Some of β_j 's will be equal to 0
- ▶ Data normalization

Prostate: Lasso

```
# loading package "lasso2"
library(lasso2)
prostate.lasso <- l1ce( lpsa ~ ., data=train, trace=TRUE, sweep.out=~1,
                        bound=seq(0,1,by=0.1) )
prostate.lasso.coef <- sapply(prostate.lasso, function(x) x$coef)
colnames(prostate.lasso.coef) <- seq( 0,1,by=0.1 )

# plot
matplot( seq(0,1,by=0.1), t(prostate.lasso.coef[-1,]), type="b",
        xlab="shrinkage factor", ylab="coefficients",
        xlim=c(0, 1.2), col="blue", pch=20 )
```



Lasso vs. Ridge regression

- ▶ Equivalent formulations

- ▶ Ridge (ℓ_2 penalty):

$$\min_{\beta} \text{RSS} \quad \text{subject to} \quad \sum_{j=1}^p \beta_j^2 \leq s$$

- ▶ Lasso (ℓ_1 penalty):

$$\min_{\beta} \text{RSS} \quad \text{subject to} \quad \sum_{j=1}^p |\beta_j| \leq s$$

- ▶ Best subset selection (ℓ_0 penalty):

$$\min_{\beta} \text{RSS} \quad \text{subject to} \quad \sum_{j=1}^p 1_{\{\beta_j \neq 0\}} \leq s$$

Last lecture: Model selection measures

Indirect measures: C_p , AIC and BIC

- ▶ Model with d predictors
- ▶ Mallows's C_p :

$$C_p = \frac{1}{n}(\text{RSS} + 2d\hat{\sigma}^2),$$

where $\hat{\sigma}$ is an estimator for the variance of noise (estimated on a model containing all predictors)

- ▶ Akaike information criteria (AIC):

$$\text{AIC} = \frac{1}{n\hat{\sigma}^2}(\text{RSS} + 2d\hat{\sigma}^2)$$

- ▶ Bayesian AIC (BIC):





$$\text{BIC} = \frac{1}{n}(\text{RSS} + d\hat{\sigma}^2 \log n)$$

- ▶ **Heuristic:** select a model with the lowest C_p , AIC, BIC, or

$$\text{Adj}R^2 = 1 - \frac{\text{RSS}/(n-d-1)}{\text{TSS}/(n-1)}$$

Example

```
> regfit.full<-regsubsets(Salary~.,data=Hitters,nvmax = 19)
> reg.summary<-summary(regfit.full)
> names(reg.summary)
[1] "which" "rssq" "rss" "adjr2" "cp" "bic" "outmat" "obj"
> reg.summary$rsq
[1] 0.3214501 0.4252237 0.4514294 0.4754067 0.4908036 0.5087146 0.5141227 0.5285569 0.5346124 0.5404950
[11] 0.5426153 0.5436302 0.5444570 0.5452164 0.5454692 0.5457656 0.5459518 0.5460945 0.5461159
```



Dimension Reduction

- ▶ Subset selection, shrinkage methods:
 - ▶ Original predictors used
- ▶ Dimensionality reduction idea
 - ▶ Represent/approximate X with a vector Z having less dimensions
 - ▶ Then, apply regression to Z
- ▶ Many approaches for doing this
- ▶ Common approach: Principal Component Analysis (PCA)

Dimension Reduction: PCA

- ▶ *Transform the predictors.* Let $Z = (Z_1, \dots, Z_q)$ represent $q < p$ linear combinations of the original p predictors:

$$Z_m = \sum_{j=1}^p \phi_{m,j} X_j$$

for some constants $\phi_{m,1}, \dots, \phi_{m,p}$

- ▶ *Use Ordinary Least Squares (OLS):* fit the linear regression

$$\hat{y}_i = \theta_0 + \sum_{m=1}^q \theta_m z_{i,m}$$

- ▶ If $\{\phi_{m,i}\}$ are chosen appropriately, dimension reduction can outperform the OLS regression

Dimension reduction

- ▶ Regression coefficients

$$\begin{aligned}\sum_{m=1}^q \theta_m z_{i,m} &= \sum_{m=1}^q \theta_m \sum_{j=1}^p \phi_{m,j} x_{i,j} \\ &= \sum_{j=1}^p \sum_{m=1}^q \theta_m \phi_{m,j} x_{i,j} \\ &= \sum_{j=1}^p \beta_j x_{i,j}\end{aligned}$$

where

$$\beta_j = \sum_{m=1}^q \theta_m \phi_{m,j}$$

- ▶ PCA is equivalent to imposing constraints on β_j 's in OLS

Finding Principal Components

PCA: Unsupervised method - will be covered more after the midterm
Good for high-dimensional data

Few words here: Finding the first principal component

- ▶ Look for the linear combination of the sample feature of the form

$$z_{i1} = \phi_{11}x_{i1} + \phi_{21}x_{i2} + \cdots + \phi_{p1}x_{ip}$$

- ▶ *First loading vector/principal component*: $\phi_1 = (\phi_{11}, \phi_{21}, \dots, \phi_{p1})$
- ▶ Assume x_i -s are **centered** ($\sum x_i = 0$)
- ▶ Look for ϕ_1 that has the largest sample variance, i.e.

$$\max_{\phi_1} \frac{1}{n} \sum_{i=1}^n z_{i1}^2 = \max_{\phi_1} \frac{1}{n} \sum_{i=1}^n (\phi_{11}x_{i1} + \phi_{21}x_{i2} + \cdots + \phi_{p1}x_{ip})^2$$

subject to $\sum_{j=1}^p \phi_{j1}^2 = 1$ (i.e., ϕ_1 is a unit vector)

- ▶ This optimization problem can be solved via eigen-decomposition

First Principal Components: Geometric interpretation

- ▶ Loading vector ϕ_1 represents the direction along which the data varies the most
- ▶ If we project x_1, \dots, x_n onto ϕ_1 , the projected values are the PC scores z_{i1} since

$$z_{i1} = \langle x_i, \phi_1 \rangle$$

Second and higher principal components

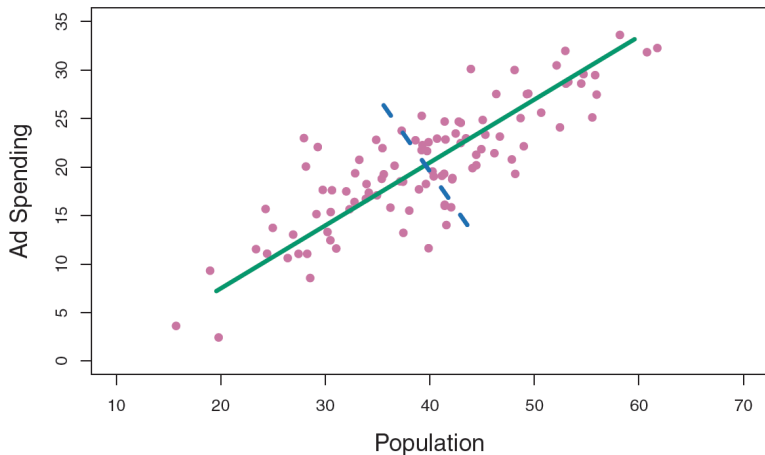
- ▶ After ϕ_1 has been determined, we look for ϕ_2 in a similar way, but with the additional constraint that ϕ_1, ϕ_2 are uncorelated

$$\langle \phi_1, \phi_2 \rangle = 0$$

- ▶ We continue this procedure until we find as many PC as we want

PCA: Example

Two PC-s: Solid line: First PC; Dashed line: Second PC



PCA as Eigenvalue-Eigenvector Decomposition

Consider finding $k \leq p$ principal components: $\phi_1, \phi_2, \dots, \phi_k$

$$U = [\phi_1 \quad \phi_2 \quad \cdots \quad \phi_k], \quad U^\top U = I_k.$$

Then, the projection of a data point $\mathbf{x}_i, 1 \leq i \leq n$ is given by

$$\hat{\mathbf{x}}_i = (\mathbf{x}_i \cdot \phi_1)\phi_1 + \cdots (\mathbf{x}_i \cdot \phi_k)\phi_k = UU^\top \mathbf{x}_i,$$

implying

$$\|\hat{\mathbf{x}}_i\|^2 = \mathbf{x}_i^\top UU^\top UU^\top \mathbf{x}_i = \mathbf{x}_i^\top UU^\top \mathbf{x}_i.$$

Note that $\hat{\mathbf{x}}_i$ minimizes the distance $\|\mathbf{x}_i - \hat{\mathbf{x}}_i\|$, and thus, finding k principle components is equivalent to finding U that maximizes

$$\begin{aligned} M &= \max_{U: U^\top U = I_k} \sum_{i=1}^n \mathbf{x}_i^\top UU^\top \mathbf{x}_i \\ &= \text{trace} \left(U^\top \sum_{i=1}^n \mathbf{x}_i \mathbf{x}_i^\top U \right), \quad (\text{using } \mathbf{x}^\top \mathbf{y} = \text{trace}(\mathbf{x} \mathbf{y}^\top)) \end{aligned}$$

PCA as Eigenvalue-Eigenvector Decomposition

Now, if $\mathbf{A} = \sum_{i=1}^n \mathbf{x}_i \mathbf{x}_i^\top$, the optimization problem becomes

$$M = \max_{\mathbf{U}: \mathbf{U}^\top \mathbf{U} = \mathbf{I}_k} \text{trace}(\mathbf{U}^\top \mathbf{A} \mathbf{U})$$

Note that \mathbf{A} is a symmetric matrix, and therefore orthogonally diagonalizable with $\lambda_1 \geq \lambda_2 \geq \dots \geq \lambda_n \geq 0$. If \mathbf{U} is composed of k eigenvectors that correspond to the k largest eigenvalues, then

$$M \geq \sum_{i=1}^k \lambda_i.$$

On the other hand, we can diagonalize $\mathbf{A} = \mathbf{V}^\top \mathbf{\Lambda} \mathbf{V}$, where \mathbf{V} is an orthogonal matrix, yielding

$$\begin{aligned} M &= \max_{\mathbf{U}: \mathbf{U}^\top \mathbf{U} = \mathbf{I}_k} \text{trace}(\mathbf{U}^\top \mathbf{V}^\top \mathbf{\Lambda} \mathbf{V} \mathbf{U}) = \max_{\mathbf{B}: \mathbf{B}^\top \mathbf{B} = \mathbf{I}_k} \text{trace}(\mathbf{B}^\top \mathbf{\Lambda} \mathbf{B}) \\ &= \sum_{i=1}^p \sum_{j=1}^k b_{ij}^2 \lambda_i = \sum_{i=1}^p \lambda_i \sum_{j=1}^k b_{ij}^2 \leq \sum_{i=1}^k \lambda_i \quad (\text{Note: } \sum_{i=1}^p b_{ij}^2 = 1, \sum_{j=1}^k b_{ij}^2 \leq 1) \end{aligned}$$

$\Rightarrow M = \sum_{i=1}^k \lambda_i$: Hence, the first k principal components correspond to the eigenvectors of the k largest eigenvalues of \mathbf{A} .

Beyond Linear - More Linear ☺: Basis Expansion

Basis expansion

- ▶ Map data \mathbf{x} into higher dimensional space $\mathbb{R}^d, d > p$:
 $\phi(\mathbf{x}) : \mathbb{R}^p \rightarrow \mathbb{R}^d$, i.e.,

$$\phi(\mathbf{x}) = (\phi_1(\mathbf{x}), \phi_2(\mathbf{x}), \dots, \phi_d(\mathbf{x}))$$

- ▶ Fit linear regression on $\phi(\mathbf{x})$ in the higher dimensional space

Example: Polynomial regression of degree q

- ▶ Map data $\mathbf{x} = (x_1, \dots, x_p)$ into

$$\phi(\mathbf{x}) = (x_1, x_1^2, \dots, x_1^q, x_2, x_2^2, \dots, x_2^q, \dots, x_p, x_p^2, \dots, x_p^q, \dots)$$

- ▶ Fit linear regression in the higher dimensional, \mathbb{R}^{pq} , space
 - ▶ Exponential growth of dimensionality: if $p = 10$ and $q = 10$

Quadratic Expansion

Example:

- ▶ Consider data with two predictors: $x_i = (x_{i1}, x_{i2})$
- ▶ We can obtain the preceding kernel by considering feature map

$$\phi(x_i) = (1, \sqrt{2}x_{i1}, \sqrt{2}x_{i2}, x_{i1}^2, \sqrt{2}x_{i1}x_{i2}, x_{i2}^2)$$

if $x \in \mathbb{R}^p$, then $\phi(x) \in \mathbb{R}^{2p}$

- ▶ Then, the inner product

$$\begin{aligned}\langle \phi(x_i), \phi(x_j) \rangle &= 1 + 2x_{i1}x_{j1} + 2x_{i1}x_{j1} + x_{i1}^2x_{j1}^2 \\ &\quad + 2x_{i1}x_{i2}x_{j1}x_{j2} + x_{i2}^2x_{j2}^2 \\ &= 1 + 2\langle x_i, x_j \rangle + \langle x_i, x_j \rangle^2 \\ &= (1 + \langle x_j, x_i \rangle)^2 = K(x_i, x_j)\end{aligned}$$

Quadratic Kernel: $K(x_i, x_j) = (1 + \langle x_j, x_i \rangle)^2$

K remains $n \times n$ matrix: does not grow with basis expansion

Polynomial Kernel and Dual Solution for Ridge

Polynomial Kernel:

$$K(x_i, x_j) = (1 + \langle x_j, x_i \rangle)^d$$

K is $n \times n$, doesn't grow with d ; recall, n is the number of data points x_i . Computing K is of the order $O(n^2p)$.

- Dual solution - recall dual form for $\hat{\mathbf{y}}$

$$\hat{\mathbf{y}} = \mathbf{K}\boldsymbol{\alpha},$$

- Dual Ridge: find $\boldsymbol{\alpha}$ that minimizes

$$\|\mathbf{y} - \hat{\mathbf{y}}\|_2^2 + \lambda \|\boldsymbol{\alpha}\|_2^2 = \|\mathbf{y} - \mathbf{K}\boldsymbol{\alpha}\|_2^2 + \lambda \|\boldsymbol{\alpha}\|_2^2$$

- Solution

$$\boldsymbol{\alpha} = (\lambda \mathbf{I} + \mathbf{K})^{-1} \mathbf{y}$$

which has computational complexity $O(n^3)$, and does not depend on the polynomial basis expansion parameter d . Hence the total complexity is $O((p+n)n^2) \approx O(n^3)$, when $p \ll n$, and does not depend on d .

Piecewise Polynomial: Splines

Example: Piecewise linear

- ▶ Assume $x \in \mathbb{R}$, i.e., $p = 1$
- ▶ Consider points: $w_1 \leq w_2 \leq \dots \leq w_n$ - called **knots**
- ▶ Consider basis

$$b_0(x) = 1, b(x) = x, b_i = (x - w_i)^+$$

- ▶ Then, fitting linear regression to these bases will result in piecewise linear approximation
- ▶ Relation to Neural Nets: Tow layer neural network corresponds to **free knot linear spline** (studied in statistics since the 70s):
 - ▶ By free knot, one means that knots/weights w_i and linear regression coefficients β_j are optimized at the same time
 - ▶ This problem is not convex

In general, **basis can be chosen to be anything**: trigonometric functions, wavelets, etc.

Where are we now?

Recall, that we started with a general supervised Statistical (Machine) Learning problems

$$Y = f(X)$$

Problem: Estimate f from training data $\{(x_i, y_i)\}$, and then use it in general for inference and prediction

- ▶ First, we assumed some characteristics of the approximation function. e.g.: \hat{f} is linear or nonlinear via basis expansion
- ▶ Then, we used the training data $\{(x_i, y_i)\}$ to find the best fit of \hat{f} to the training data by minimizing the square error, RSS, or some other error function.
- ▶ For linear regression and Gaussian noise, we used classical tools from statistics, χ^2 , t-statistics and F-statistics, to characterize the confidence of our model.

Where are we now and what are we doing next?

- ▶ Searching for simpler models:
 - ▶ Selecting a model using t -statistics and F -statistics
 - ▶ Trying all combinations of predictors - not scalable - and then using indirect measures (C_p , AIC, BIC, etc.).
 - ▶ Introducing penalties: Ridge or Lasso - scalable
 - ▶ Preprocessing data - dimensionality reduction
- ▶ Bias - Variance: in general the preceding procedures introduce bias, but the hope is that at an expense of a small bias we have a bigger reduction in variance
- ▶ How do we select the best model?
 - ▶ For linear and Gaussian: we could use the t -statistics and F -statistics
 - ▶ We could use the indirect measures: C_p , AIC, BIC, etc.
 - ▶ Are there more direct general procedures to test and select the best model?

Model validation: Training vs. test error

- ▶ Select a statistical learning method, e.g.: linear model
- ▶ Training error: the average error from using the method to predict the response on the observations used in its training
- ▶ **Test error**: the average error from using the method to predict the response on a **new observation**
- ▶ Ideally: a large designated test set – seldomly available

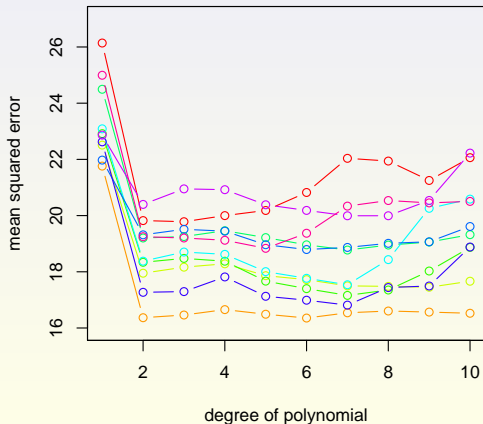
Validation-set approach

- ▶ Randomly divide the available samples into:
 - ▶ training set
 - ▶ validation set
- ▶ Random split into two halves
- ▶ Fit a model using the training set
- ▶ Use the model to predict the responses in the validation set
- ▶ The validation-set error is an estimate of the test error
- ▶ Drawbacks
 - ▶ The error estimate can be variable – depends on the split
 - ▶ Only a subset of observations used to fit the model
 - ▶ Tends to overestimate the test error for the model fit on the entire data set

Example: Auto data set

```
> set.seed(1)
> train<-sample(392,196) #pick randomly 1/2 sample
> err<-rep(0,10)
> for(i in 1:10) {
+   lm.fit<-lm(mpg~poly(horsepower,i),data = auto, subset = train)
+   err[i]<-mean((auto$mpg-predict(lm.fit,auto))[-train]^2)}
```

Different "train" set produces different results



K -fold cross-validation

- ▶ Popular approach
- ▶ Used in model selection
- ▶ Procedure
 - ▶ Randomly divide observations into K equal-sized parts
 - ▶ Leave out part k , fit a model using the remaining $K - 1$ parts
 - ▶ Use the left-out part to estimate the error
 - ▶ Repeat for all k
 - ▶ Combine results

K -fold cross-validation

- ▶ K parts: C_1, C_2, \dots, C_K
- ▶ $\cup_k C_k = \{1, \dots, n\}$
- ▶ n_k : the number of observations in part k
- ▶ Compute

$$CV_{(K)} = \sum_{k=1}^K \frac{n_k}{n} \text{MSE}_k$$

where

$$\text{MSE}_k = \frac{1}{n_k} \sum_{i \in C_k} (y_i - \hat{y}_i)^2$$

and \hat{y}_i is the prediction for observation i obtained from the data without part k

- ▶ $K = n$: leave-one out cross-validation (LOOCV)

LOOCV

Linear model example: we can compute CV error

- ▶ **No randomness** – all subsets of size $(n - 1)$ considered
- ▶ \mathbf{X} and \mathbf{y} :
 - ▶ observation i : \mathbf{X}_i and y_i
 - ▶ no observation i : $\mathbf{X}_{(i)}$ and $\mathbf{y}_{(i)}$
- ▶ Coefficients
 - ▶ observation i omitted:

$$\mathbf{X}_{(i)}^\top \mathbf{X}_{(i)} \hat{\boldsymbol{\beta}}_{(i)} = \mathbf{X}_{(i)}^\top \mathbf{y}_{(i)}$$

- ▶ all observations used:

$$(\mathbf{X}_{(i)}^\top \mathbf{X}_{(i)} + \mathbf{X}_i^\top \mathbf{X}_i) \hat{\boldsymbol{\beta}} = \mathbf{X}_{(i)}^\top \mathbf{y}_{(i)} + \mathbf{X}_i^\top y_i$$

and

$$\hat{\boldsymbol{\beta}} = \hat{\boldsymbol{\beta}}_{(i)} + (\mathbf{X}_{(i)}^\top \mathbf{X}_{(i)})^{-1} \mathbf{X}_i^\top (y_i - \hat{y}_i),$$

where \hat{y}_i is the prediction for y_i using all observations

$$\hat{y}_i = \mathbf{X}_i \hat{\boldsymbol{\beta}} = \mathbf{X}_i (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}$$

LOOCV

- ▶ Some algebra and

$$\begin{aligned}(\mathbf{X}_{(i)}^\top \mathbf{X}_{(i)})^{-1} &= (\mathbf{X}^\top \mathbf{X} - \mathbf{X}_i^\top \mathbf{X}_i)^{-1} \\&= (\mathbf{X}^\top \mathbf{X})^{-1} + \frac{(\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}_i^\top \mathbf{X}_i (\mathbf{X}^\top \mathbf{X})^{-1}}{1 - \mathbf{X}_i (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}_i^\top}\end{aligned}$$

yield

$$\begin{aligned}\text{CV}_{(n)} &= \frac{1}{n} \sum_{i=1}^n (y_i - \mathbf{X}_i \hat{\boldsymbol{\beta}}_{(i)})^2 \\&= \frac{1}{n} \sum_{i=1}^n \left(\frac{y_i - \hat{y}_i}{1 - h_i} \right)^2\end{aligned}$$

where $h_i = \mathbf{X}_i (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}_i^\top$


- ▶ **Weighted sum of squared residuals** (Provides validation for using modified RSS for some of the indirect measures).

Example: Auto data set


```
> library(boot)
> cv.err<-rep(0,10)
> for(i in 1:10) {
+   glm.fit<-glm(mpg~poly(horsepower,i), data=auto)
+   cv.err[i]<-cv.glm(auto,glm.fit)$delta[1]
+ }

> set.seed(1)
> for(i in 1:10) {
+   glm.fit<-glm(mpg~poly(horsepower,i), data=auto)
+   cv.err[i]<-cv.glm(auto,glm.fit,K=10)$delta[1]
+ }
```

LOOCV



10-fold CV



What if we are interested in computing other statistics, not just the test error, for the non-Gaussian or nonlinear case?

Bootstrap

- ▶ Setup

- ▶ Population model that produces an outcome Y
- ▶ Observations \mathbf{Z} from this population model
- ▶ Statistic $T(\mathbf{Z})$
- ▶ Distribution of $T(\mathbf{Z})$

- ▶ Idea

- ▶ The distribution of $T(\mathbf{Z})$ can be estimated by sampling \mathbf{Z} from the population model
- ▶ Resample with replacement from \mathbf{Z} to “approximate” sampling from the population model

- ▶ Why?

- ▶ Only samples \mathbf{Z} available
- ▶ No information on the population model

Bootstrap: Basic algorithm

- ▶ Input

- ▶ A sample of data $\mathbf{Z} = (\mathbf{Z}_1, \dots, \mathbf{Z}_n)$
- ▶ An estimation rule \hat{T} for Statistic T

- ▶ Algorithm

1. Generate bootstrap samples $\mathbf{Z}^{*1}, \mathbf{Z}^{*2}, \dots, \mathbf{Z}^{*B}$
 - ▶ Create \mathbf{Z}^{*b} by selecting n points from \mathbf{Z}
 - ▶ A particular \mathbf{Z}_i can appear in \mathbf{Z}^{*b} multiple times
2. Evaluate the estimator on each \mathbf{Z}^{*b} :

$$\hat{T}_b = \hat{T}(\mathbf{Z}^{*b})$$

- ▶ The empirical distribution of $\{\hat{T}_1, \dots, \hat{T}_B\}$ is an estimate of the distribution of $T(\mathbf{Z})$
- ▶ Bootstrap distribution
- ▶ Overlap between \mathbf{Z} and \mathbf{Z}^{*b} ?

Example: Variance estimation of the median

- ▶ The median of x_1, \dots, x_n , $x_i \in \mathbb{R}$, is found by sorting the numbers and taking the middle one (or averaging the two middle ones)
- ▶ How good is the estimate $\text{median}(x_1, \dots, x_n)$?
 - ▶ Find it's variance
 - ▶ How?

1. Generate bootstrap datasets $\mathbf{Z}^{*1}, \dots, \mathbf{Z}^{*B}$

2. Calculate:

$$\hat{T}_{\text{mean}} = \frac{1}{B} \sum_{b=1}^B \text{median}(\mathbf{Z}^{*b})$$

and



$$\hat{T}_{\text{var}} = \frac{1}{B-1} \sum_{b=1}^B \left(\text{median}(\mathbf{Z}^{*b}) - \hat{T}_{\text{mean}} \right)^2$$

When it works?

- ▶ The bootstrap distribution should be close to the sampling distribution
 - ▶ The number of bootstrap samples B should be large enough
 - ▶ The original data sample Z should be large enough to be “representative” of the population model
- ▶ Few assumptions about the population model
- ▶ Can yield inaccurate results (e.g., “extreme value” statistics)
- ▶ Usually reliable for estimating standard errors for estimators
 - ▶ Standard errors
 - ▶ The bootstrap estimate of the standard error is the standard deviation of the bootstrap distribution
 - ▶ Confidence intervals
 - ▶ $(1 - \gamma)$ confidence interval
 - ▶ t_α – the α th quantile of the bootstrap distribution
 - ▶ $(1 - \gamma)$ percentile bootstrap interval: $[t_{\gamma/2}, t_{1-\gamma/2}]$



Example: Mean

```
> library(stats)
> set.seed(100)
> z<-rexp(100,rate=1)
> boot.fn=function(data,index)
+   return(mean(data[index]))
> boot.fn(z,1:100)
[1] 0.9874761
> boot.fn(z,sample(100,100,replace=T))
[1] 1.184723
> boot.fn(z,sample(100,100,replace=T))
[1] 1.131222
> library(boot)
> boot.out<-boot(z,boot.fn,1000)
```



Example: Median

```
> boot.fn=function(data,index)
+   return(median(data[index]))
> boot.fn(z,1:100)
[1] 0.6955635
> boot.fn(z,sample(100,100,replace=T))
[1] 0.6373702
> boot.fn(z,sample(100,100,replace=T))
[1] 0.652271
> boot.out<-boot(z,boot.fn,1000)
```



Bootstrap: Regression modeling

- ▶ n observations, response \mathbf{y} , covariates \mathbf{X}
- ▶ Bootstrap standard errors for OLS coefficients using case resampling:
 - ▶ For $b = 1, \dots, B$
 - ▶ Draw sample uniformly at random, with replacement, from observations (\mathbf{X}, \mathbf{y}) . Let the i th outcome in the b th sample be $(\mathbf{X}_i^{*b}, y_i^{*b})$
 - ▶ Compute $\hat{\beta}^{*b}$ given $(\mathbf{X}^{*b}, \mathbf{y}^{*b})$
- ▶ Bootstrap distribution of $\hat{\beta}$ to compute standard errors

Example: When analytics is unavailable, use bootstrap

Example: Heteroskedasticity

- ▶ Suppose $X_i \sim \mathcal{N}(0, 1)$ and $y_i = X_i + \epsilon_i$, for $i = 1, \dots, n$, where $\epsilon_i \sim \mathcal{N}(0, X_i^4)$
- ▶ Non-constant error variance
- ▶ Standard assumptions violated
- ▶ Standard errors?

Example: Heteroskedasticity

```
> n<-1000
> set.seed(1)
> x<-rnorm(n); y<-x+x^2*rnorm(n); df<-data.frame(x,y)
>
> lm.fit<-summary(lm(y~x,data = df))
> lm.fit$coefficients
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	0.02807759	0.06502742	0.4317807	6.659940e-01
x	1.18461640	0.06286099	18.8450180	5.229894e-68

```
>
> coef.boot=function(data,indices) {
+   coef.fit<-lm(y~x,data=data[indices,])
+   return(coef(coef.fit))
+ }
>
> boot.out<-boot(df,coef.boot, 50000)
> boot.out
```

ORDINARY NONPARAMETRIC BOOTSTRAP

Call:

```
boot(data = df, statistic = coef.boot, R = 50000)
```

Bootstrap Statistics :

	original	bias	std. error
t1*	0.02807759	-0.0008122518	0.06504296
t2*	1.18461640	-0.0014536288	0.14104734

Reading:

ISL: Read Chapters 5 and 7

ESL: Read Chapter 5, 7, and specifically Sections 5.1-5.2 and 7.10-7.11 for this lecture.

Homework 1: Due today, Sep 27, by 11:59pm.

Homework 2: Due Fri, Oct 7th, by 11:59pm.

Midterm planned for Oct 25th